



US007747601B2

(12) **United States Patent**
Cooper et al.

(10) **Patent No.:** **US 7,747,601 B2**
(45) **Date of Patent:** **Jun. 29, 2010**

(54) **METHOD AND APPARATUS FOR IDENTIFYING AND CLASSIFYING QUERY INTENT**

(75) Inventors: **Edwin Riley Cooper**, Oakland, CA (US); **Michael Peter Dukes**, Sunnyvale, CA (US); **Gann Alexander Bierner**, Oakland, CA (US); **Filippo Ferdinando Paulo Beghelli**, Reading (GB)

(73) Assignee: **Inqura, Inc.**, San Bruno, CA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 77 days.

(21) Appl. No.: **11/464,443**

(22) Filed: **Aug. 14, 2006**

(65) **Prior Publication Data**

US 2009/0077047 A1 Mar. 19, 2009

(51) **Int. Cl.**
G06F 7/00 (2006.01)
G06F 17/30 (2006.01)

(52) **U.S. Cl.** **707/708**

(58) **Field of Classification Search** None
See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

5,321,833 A	6/1994	Chang et al.	
5,386,556 A	1/1995	Hedin et al.	
5,535,382 A	7/1996	Ogawa	
5,608,624 A	3/1997	Luciw	
5,625,814 A	4/1997	Luciw	
5,694,523 A *	12/1997	Wical	706/45
5,694,546 A	12/1997	Reisman	
5,742,816 A	4/1998	Barr et al.	
5,794,050 A	8/1998	Dahlgren et al.	
5,812,865 A	9/1998	Theimer et al.	
5,826,269 A	10/1998	Hussey	
5,848,399 A	12/1998	Burke	

5,873,056 A	2/1999	Liddy et al.
5,873,076 A	2/1999	Barr et al.
5,873,080 A	2/1999	Coden et al.
5,878,423 A	3/1999	Anderson et al.
5,884,302 A	3/1999	Ho
5,890,152 A	3/1999	Rapaport et al.
5,893,091 A	4/1999	Hunt et al.

(Continued)

FOREIGN PATENT DOCUMENTS

WO WO 0188662 A2 11/2001

(Continued)

OTHER PUBLICATIONS

Madhavan et al, "Semantic Mappings for Data Mediation", Feb. 2002, <pages.cs.wisc.edu/~anhai/talks/JayantMadhavan-Affiliates2002.ppt>, pp. 1-17.

(Continued)


Primary Examiner—Mohammad Ali
Assistant Examiner—Alexey Shmatov

(74) *Attorney, Agent, or Firm*—Stolowitz Ford Cowger LLP

(57) **ABSTRACT**

Linguistic analysis is used to identify queries that use different natural language formations to request similar information. Common intent categories are identified for the queries requesting similar information. Intent responses can then be provided that are associated with the identified intent categories. An intent management tool can be used for identifying new intent categories, identifying obsolete intent categories, or refining existing intent categories.

18 Claims, 16 Drawing Sheets



U.S. PATENT DOCUMENTS

5,897,622 A 4/1999 Blinn et al.
 5,901,287 A 5/1999 Bull et al.
 5,913,215 A 6/1999 Rubinstein et al.
 5,933,822 A 8/1999 Braden-Harder et al.
 5,948,054 A 9/1999 Nielsen
 5,963,940 A * 10/1999 Liddy et al. 707/5
 5,966,695 A 10/1999 Melchione et al.
 5,974,412 A 10/1999 Hazlehurst et al.
 5,987,454 A 11/1999 Hobbs
 5,995,921 A 11/1999 Richards et al.
 6,006,225 A 12/1999 Bowman et al.
 6,016,476 A 1/2000 Maes et al.
 6,021,403 A 2/2000 Horvitz
 6,026,388 A 2/2000 Liddy et al.
 6,028,601 A 2/2000 Machiraju et al.
 6,052,710 A 4/2000 Saliba et al.
 6,061,057 A 5/2000 Knowlton et al.
 6,070,149 A 5/2000 Tavor et al.
 6,078,914 A 6/2000 Redfern
 6,151,600 A 11/2000 Dedrick
 6,208,991 B1 3/2001 French et al.
 6,233,547 B1 5/2001 Denber
 6,278,996 B1 8/2001 Richardson et al.
 6,282,534 B1 8/2001 Vora
 6,285,998 B1 9/2001 Black et al.
 6,370,535 B1 4/2002 Shapiro
 6,385,592 B1 5/2002 Angles et al.
 6,393,479 B1 5/2002 Glommen et al.
 6,401,084 B1 * 6/2002 Ortega et al. 707/2
 6,401,094 B1 6/2002 Stemp et al.
 6,466,899 B1 10/2002 Yano
 6,471,521 B1 10/2002 Dornbush et al.
 6,477,551 B1 11/2002 Johnson et al.
 6,480,843 B2 11/2002 Li
 6,584,464 B1 6/2003 Warthen
 6,675,159 B1 1/2004 Lin et al.
 6,766,320 B1 * 7/2004 Wang et al. 707/5
 6,907,414 B1 6/2005 Parnell
 6,928,425 B2 8/2005 Grefenstette et al.
 6,941,301 B2 9/2005 Fergusson et al.
 6,957,213 B1 10/2005 Yuret
 7,024,400 B2 4/2006 Tokuda et al.
 7,035,864 B1 4/2006 Ferrari et al.
 7,177,795 B1 2/2007 Chen et al.
 7,428,541 B2 9/2008 Houle
 2002/0051020 A1 5/2002 Ferrari et al.
 2002/0078090 A1 6/2002 Hwang et al.
 2002/0103809 A1 8/2002 Starzl et al.
 2003/0014403 A1 1/2003 Chandrasekar et al.
 2003/0037073 A1 2/2003 Tokuda et al.
 2003/0144994 A1 * 7/2003 Wen et al. 707/3
 2003/0204435 A1 10/2003 McQuilkin et al.
 2003/0217052 A1 11/2003 Rubenczyk et al.
 2003/0233224 A1 12/2003 Marchisio et al.
 2004/0024739 A1 2/2004 Copperman et al.
 2004/0167870 A1 8/2004 Wakefield et al.
 2005/0080775 A1 4/2005 Colledge et al.

2005/0203878 A1 9/2005 Brill
 2005/0267871 A1 12/2005 Marchisio et al.
 2006/0074836 A1 4/2006 Gardner
 2006/0106769 A1 5/2006 Gibbs
 2006/0122979 A1 * 6/2006 Kapur et al. 707/3
 2006/0136403 A1 6/2006 Koo
 2006/0179404 A1 8/2006 Yolleck et al.
 2006/0253427 A1 11/2006 Wu
 2008/0294651 A1 11/2008 Masuyama et al.

FOREIGN PATENT DOCUMENTS

WO WO 0235376 A2 5/2002
 WO WO/2007/134128 11/2007
 WO WO/2008/067316 5/2008

OTHER PUBLICATIONS

Chien, Steve, et al. "Semantic similarity between search engine queries using temporal correlation." In Proceedings of the 14th international conference on World Wide Web, Chiba, Japan. Session: Usage analysis, p. 2-11. Published May 10-14, 2005. [retrieved Jan. 21, 2008]. Retrieved from the Internet: <URL: <http://www.ra.ethz.ch/CDstore/www2005/docs/p2.pdf>>.
 Raghavan, S. et al. Crawling the Hidden Web. Computer Science Department, Stanford University, 2001, pp. 1-25, especially pp. 1-15.
 Don Clark, "AnswerFriend Seeks to Sell Question-Answering Software," Wall Street Journal, Aug. 24, 2000.
 Deniz Yuret, "Discovery of Linguistic Relations Using Lexical Attraction" PhD Thesis, MIT, May 15, 1998.
 Deniz Yuret, "Lexical Attraction Models of Language." Submitted to The Sixteenth National Conference on Artificial Intelligence, 1999.
 Boris Katz, Deniz Yuret, et al. "Integrating Large Lexicons and Web Resources into a Natural Language Query System." To appear in Proceedings of the IEEE International Conference on Multimedia Computing and Systems, 1999.
 International Preliminary Report on Patentability and Written Opinion of the International Searching Authority for PCT/US2007/068597; International Searching Authority/US; Nov. 11, 2008.
 International Search Report for PCT/US2007/068597; International Searching Authority/US; May 22, 2008.
 International Preliminary Report on Patentability and Written Opinion of the International Searching Authority for PCT/US2007/075929; International Searching Authority/US; Feb. 17, 2009.
 International Search Report for PCT/US2007/075929; International Searching Authority/US; Jan. 21, 2008.
 International Preliminary Report on Patentability and Written Opinion of the International Searching Authority for PCT/US2007/085646; International Searching Authority/US; Jun. 3, 2009.
 International Search Report for PCT/US2007/085646; International Searching Authority/US; Jul. 23, 2008.
 International Search Report for PCT/US2001/15711; International Searching Authority/US; Nov. 6, 2001.
 International Preliminary Examination Report for PCT/US2001/15711; IPEA/US; Date of completion of report Feb. 16, 2005.
 Stolzowicz Ford Cowger LLP, Listing of Related Cases, Aug. 19, 2009. Internet Archive, Wayback Machine, website pages from LLBean, Feb. 29, 2000.

* cited by examiner




FIG.1




FIG. 2





FIG.3A






FIG.4

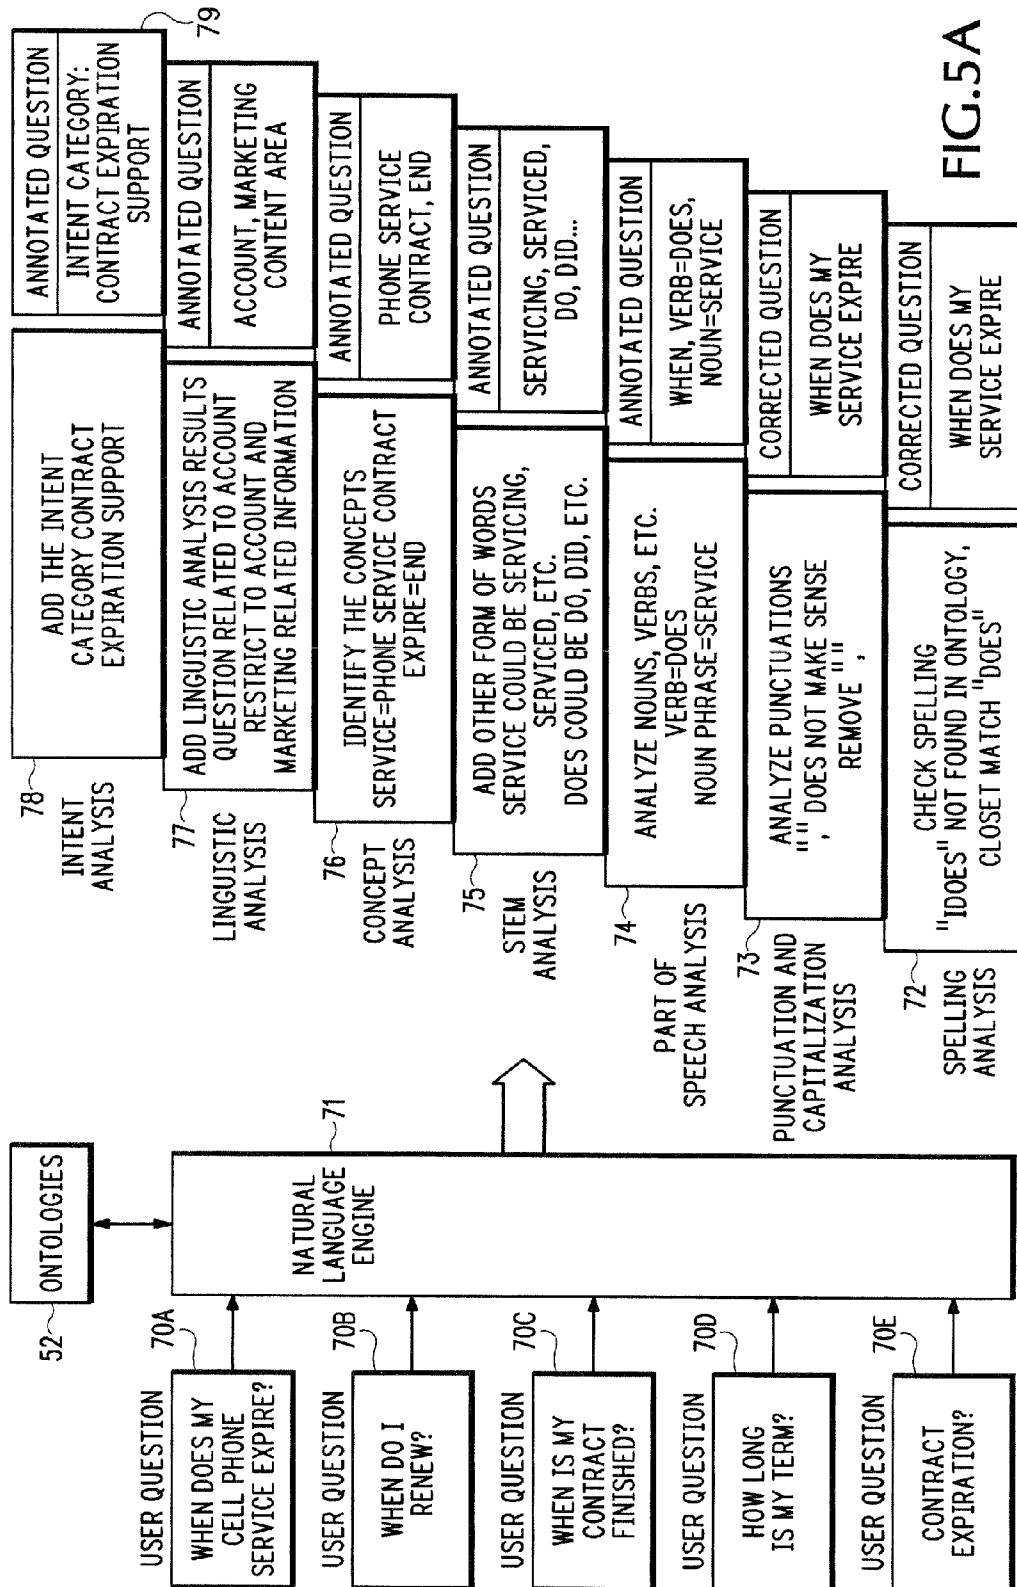

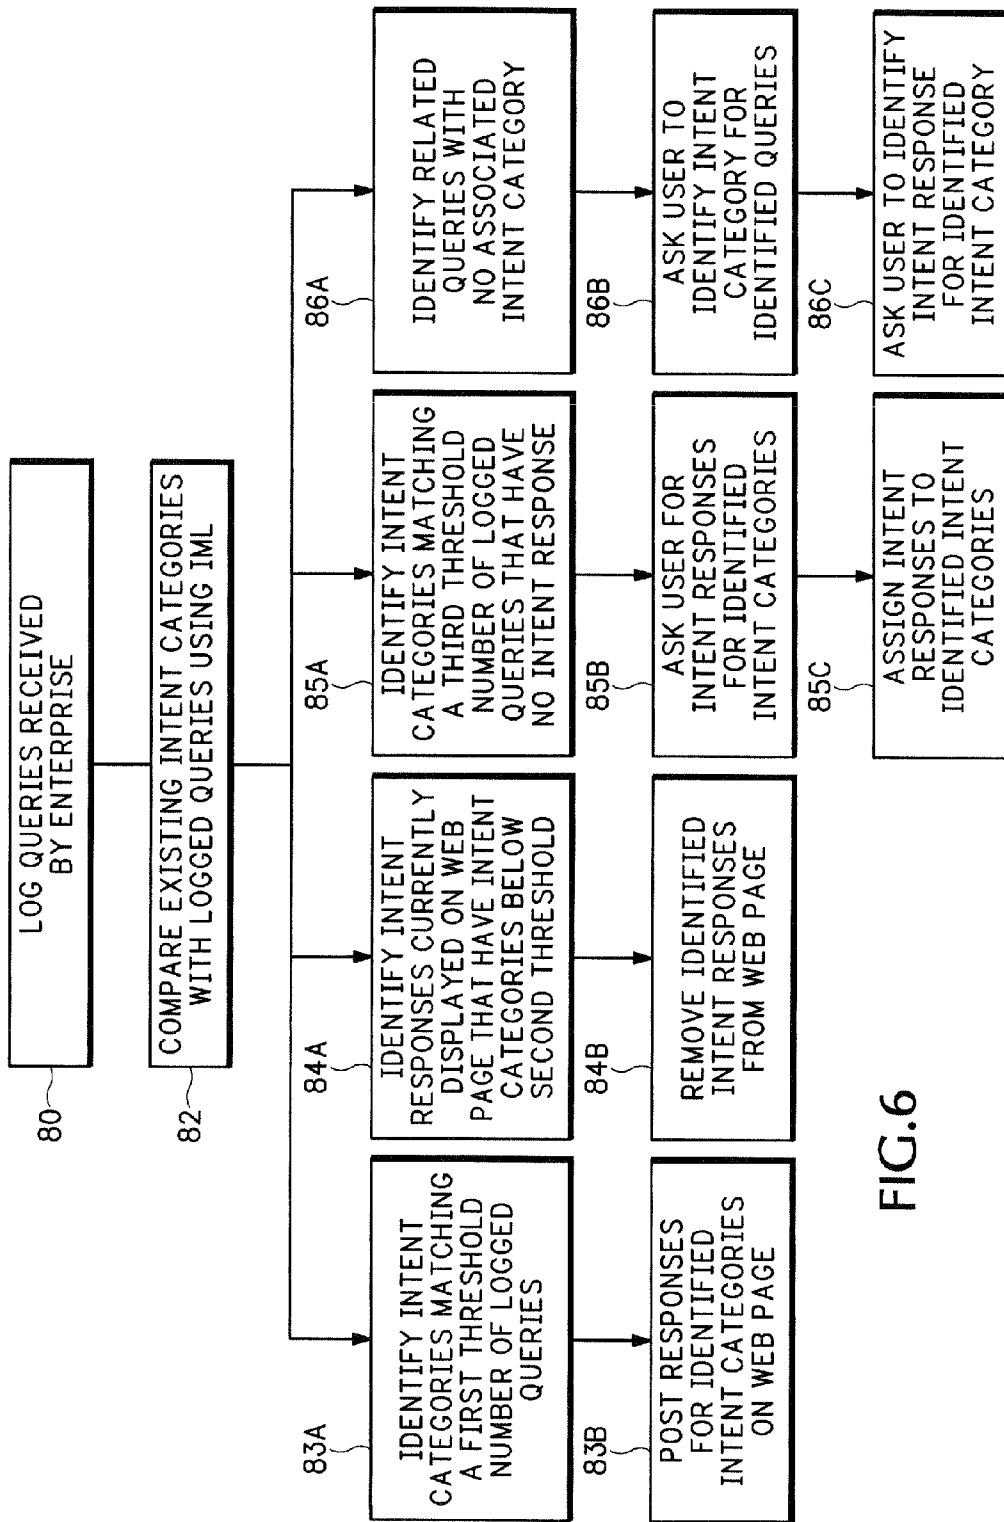


FIG. 5A





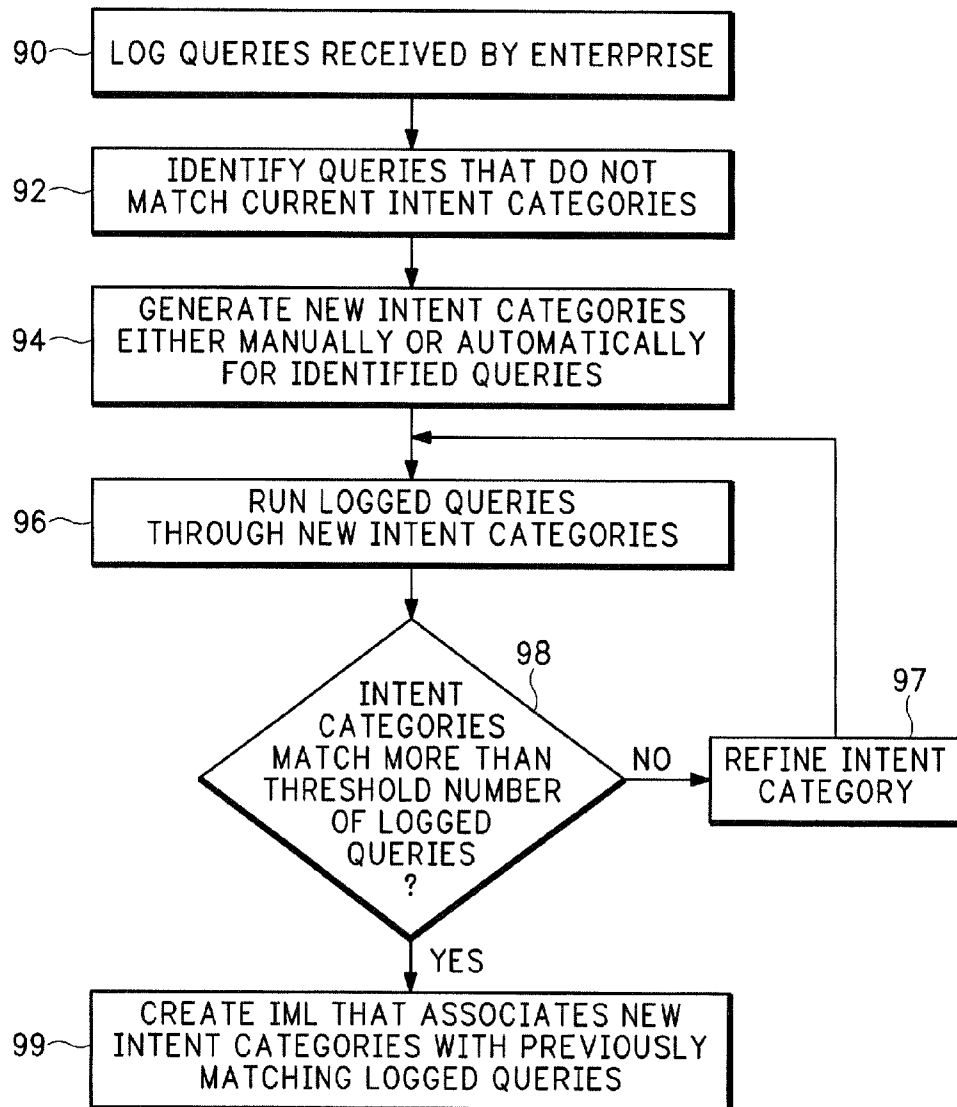


FIG.7

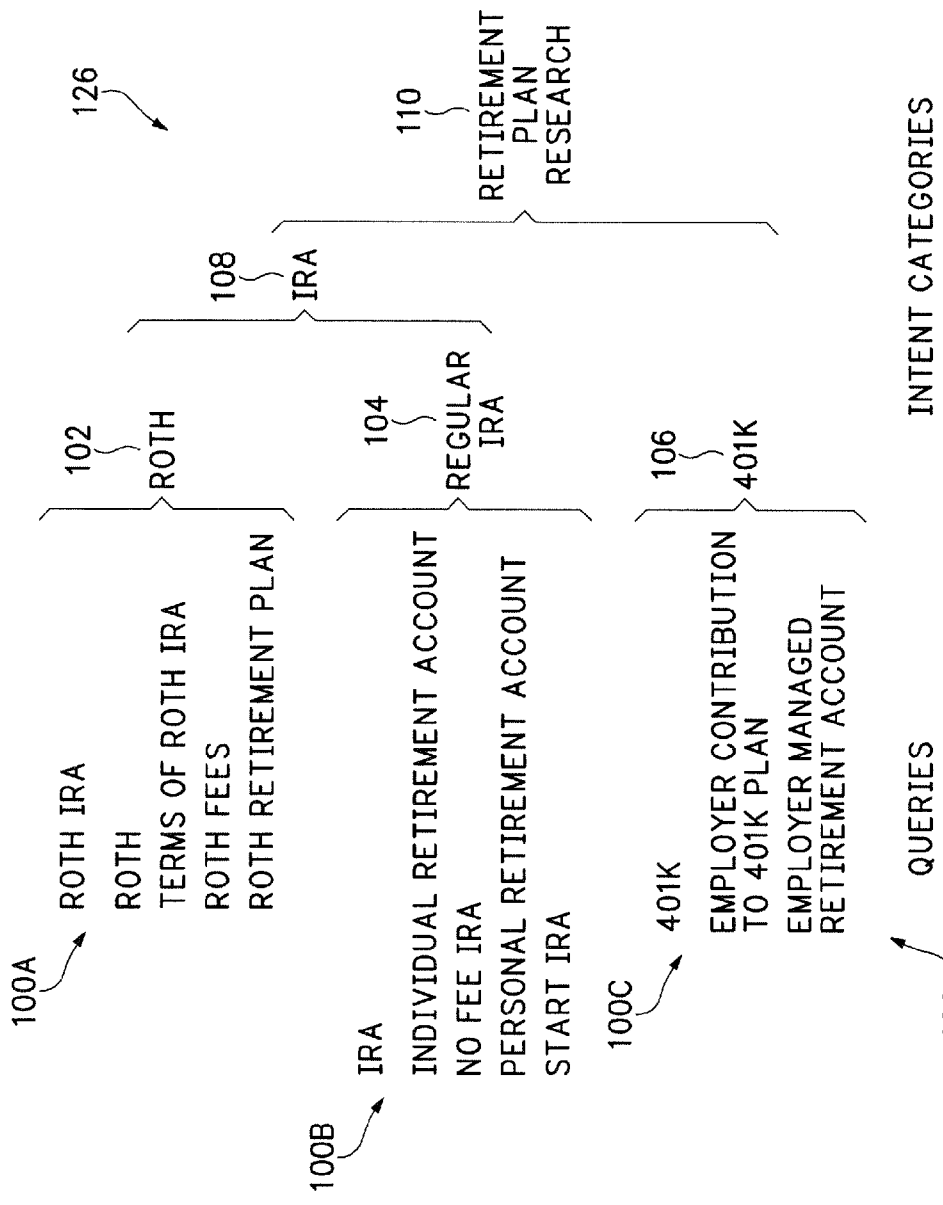


FIG.8





FIG.9



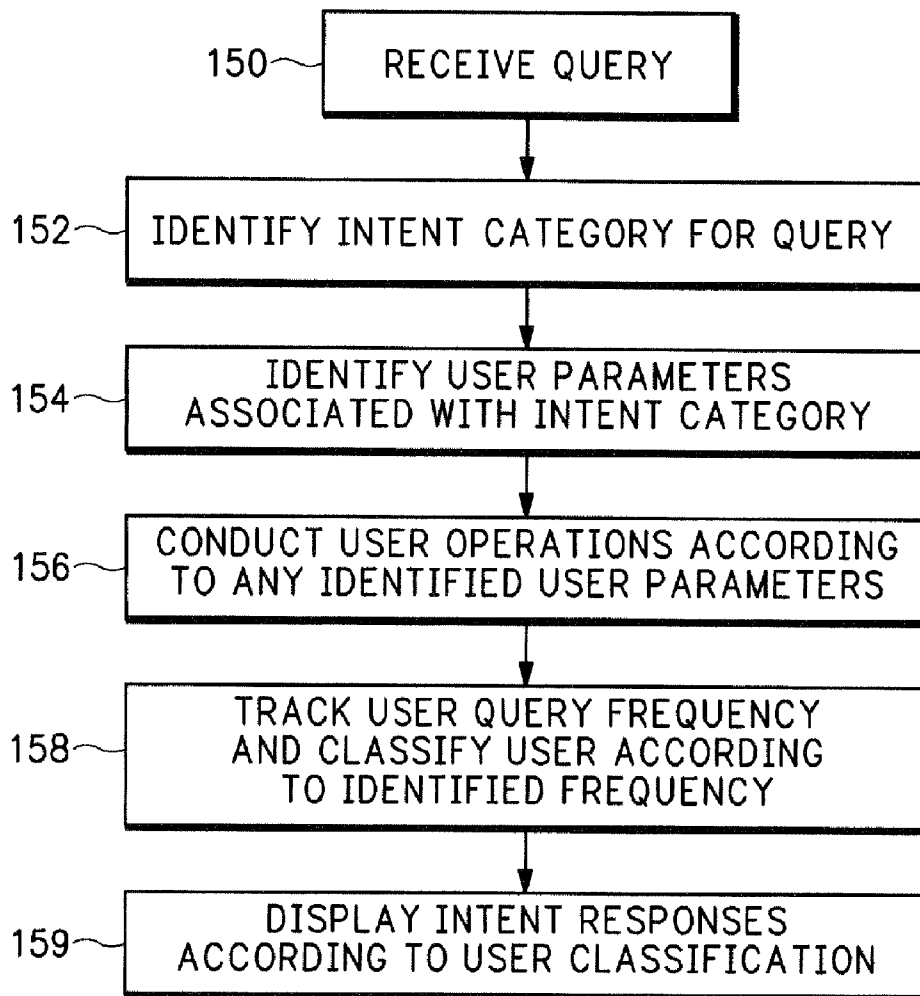


FIG.11

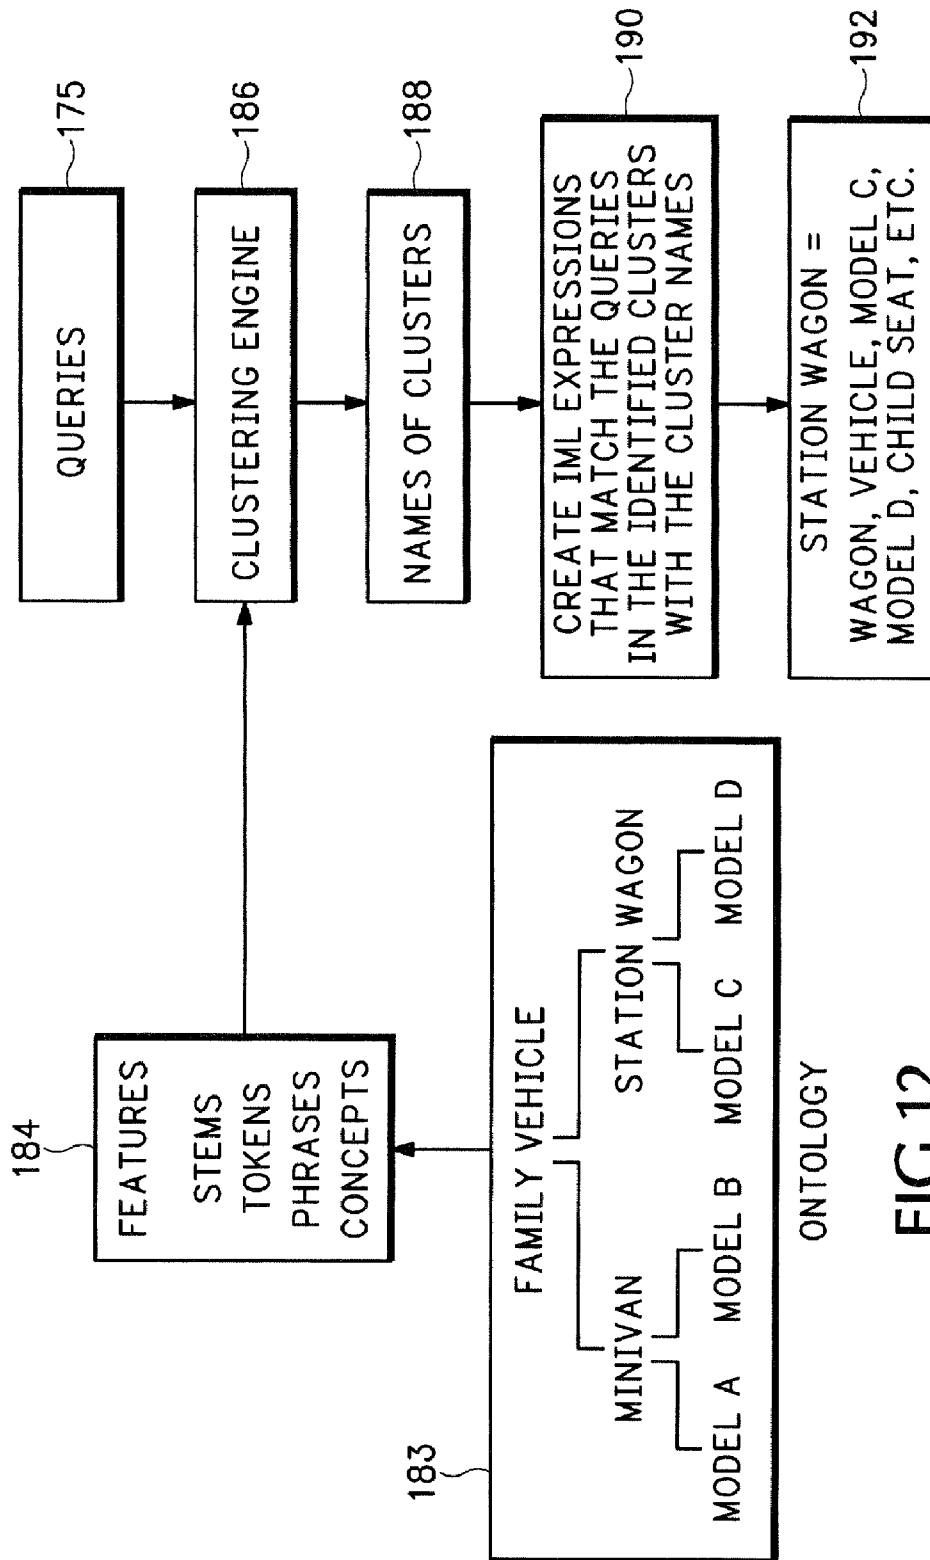


FIG.12


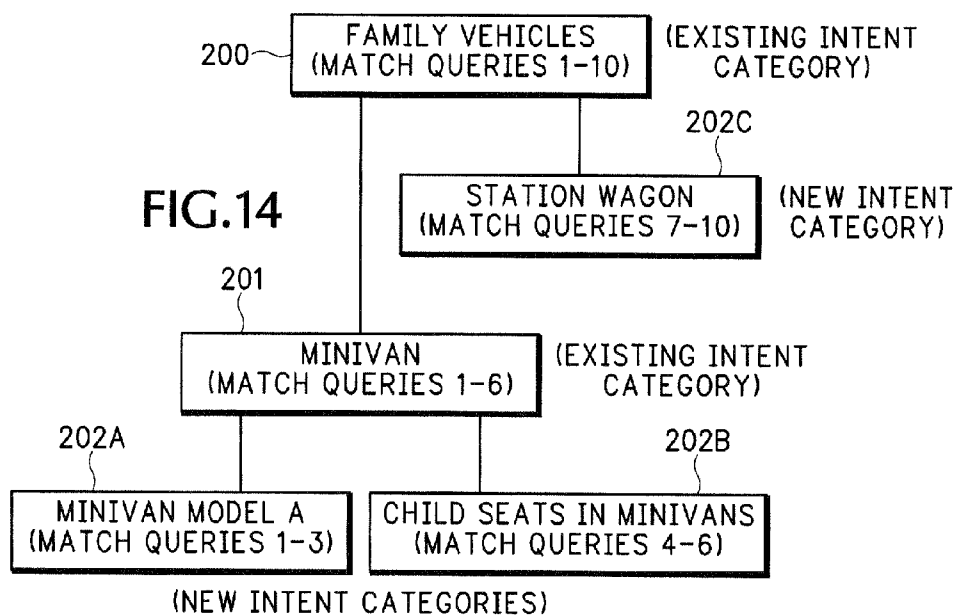


FIG.13






FIG.15

1

METHOD AND APPARATUS FOR IDENTIFYING AND CLASSIFYING QUERY INTENT

BACKGROUND

Search engines try to provide the most relevant responses to user questions. Unfortunately, many search engines return information that may be unrelated, or not directly related, to the question. For example, search engines may return any document containing words matching keywords in the question. The user has to then manually sort through each returned document in an attempt to identify information that may be relevant or answer the question. This “brute force” method is time consuming and often fails to locate the precise information sought in the question.

Current search engines try to help the user in their manual document search by ranking the returned documents. This ranking method may rank documents simply according to the number of words in the documents that match keywords in the query. At least one critical limitation with this keyword search technique is that the user may not necessarily input search terms needed by the search engine to locate the correct information. In addition, even appropriate keywords may also be associated with other documents unrelated to the information sought by the user.

Search engines have been developed that attempt to classify queries. For example, the search engine may try to associate different words in the search query with different information categories. The search engine then attempts to provide the user with responses associated with the identified information category.

A critical problem with these information retrieval schemes is that there are seemingly limitless ways in a natural language for a user to request for the same information. And as also mentioned above, the user may not necessarily enter, or even know, the best words or phrases for identifying the appropriate information. Accordingly, the search engine can only classify a very limited number of subject matters. Further, a large amount of human resources are required to keep these types of search engines up to date with new information categories that may develop over time. Thus, these “higher level” search engines have had only limited success providing responses to user questions.

The present invention addresses this and other problems associated with the prior art.

SUMMARY OF THE INVENTION

Linguistic analysis is used to identify queries that use different natural language formations to request similar information. Common intent categories are identified for the queries requesting similar information. Intent responses can then be provided that are associated with the identified intent categories. An intent management tool can be used for identifying new intent categories, identifying obsolete intent categories, or refining existing intent categories.

The foregoing and other objects, features and advantages of the invention will become more readily apparent from the following detailed description of a preferred embodiment of the invention which proceeds with reference to the accompanying drawings.

BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 is a graph that shows how the number of unique queries received by an enterprise can be reduced by classifying the queries into intent categories.

2

FIG. 2 is a block diagram showing an intent based search engine.

FIG. 3A is a graph showing how the intent based search engine can provide different types of responses according to query frequency.

FIG. 3B is a block diagram showing the intent based query engine in more detail.

FIG. 4 is a block diagram showing how an intent management tool is used for managing intent categories.

FIG. 5A is a block diagram showing how intent categories are automatically identified and derived from queries.

FIG. 5B shows in more detail how intent categories are automatically identified and derived from queries.

FIG. 6 is a flow diagram showing in more detail how the intent management tool in FIG. 4 manages intent categories.

FIG. 7 is a flow diagram showing how new intent categories can be generated and refined using the intent management tool.

FIG. 8 shows how intent categories can be associated with an intent hierarchy.

FIG. 9 is a block diagram showing how different intent responses are displayed according to an associated intent hierarchy.

FIG. 10 is a block diagram showing how a user or administrator can associate different parameters with an intent category.

FIG. 11 is a flow diagram showing how user parameters can be associated with an intent category.

FIG. 12 is a block diagram showing how features from an ontology are used to identify query clusters.

FIG. 13 is a flow diagram showing how clustering is used to generate new intent categories.

FIG. 14 is a block diagram showing how new intent categories identified in FIG. 13 are associated with different locations in an intent hierarchy.

FIG. 15 is a block diagram showing how parameters can be assigned to intent responses.

DETAILED DESCRIPTION

FIG. 1 is a graph showing the relationship between unique search queries and the frequency they are received by a particular enterprise. The enterprise could be any type of business or institution, such as a car dealer, financial service, telecommunications company, etc. that has a search engine that attempts to provide responses to questions submitted by users via computer terminals. The horizontal axis 12 refers to the different unique queries that may be received by one or more search engines operated by the enterprise. The vertical axis 14 refers to the frequency of the different unique queries.

A challenge exists trying to algorithmically decipher the meaning of received queries and then provide responses associated with the identified query meaning. For example, the shape of curve 16 indicates that a first portion 20 of the unique queries occur with the most frequency and a second portion 18 of the queries occur with slightly less frequency. As can be seen, a large percentage of the total number of queries occur in this second lower frequency portion 18.

Due to maintenance and resource issues, it may only be possible for search engines to try and determine the meaning and then provide associated responses for a subset of the most frequency received queries. For example, a search engine for an online book retailer may be designed to look for and identify queries related to book prices. However, it may not be cost and time effective to design the search engine to try and determine the meaning and provide associated responses for every possibly book question. For example, thousands of

possible responses would have to be configured just to capture a relatively small percentage of possible book questions. This large number of preconfigured responses are difficult to maintain and would have to be constantly updated to respond to the never ending number of new questions related to new books.

Unfortunately and according to curve 16, developing a search engine that is only capable of responding to the most frequency asked questions 20, ignores a large percentage of queries 18 that may be received by the online book retailer. This substantial portion of "query outliers" 18 would then have to be processed using conventional keyword searches. The limitations of key word searching was previously explained above.

As a result, the search engine for the online book retailer may not provide the most relevant responses for a large percentage of the received queries. This could negatively effect business. In the online book seller example, some of the less frequently received queries 18 may relate to rare books that may have larger mark ups than the more recent/popular books associated with the most frequently received queries 20. Accordingly, these search engine limitations may cause the online book retailer to lose some high profit rare books sales.

An intent based search engine is used to determine the intent categories of queries and then provide corresponding responses for a larger proportion of unique queries that may be received by an enterprise. "Intent" refers to the meaning associated with a query. An intent based search engine classifies multiple different unique queries into common "useful" intent categories. The term "useful" refers to intent categories that are associated with relevant responses to information requests. For example, identifying an intent category for a group of queries associated with "the Internet", may be too broad to be useful to an enterprise that is attempting to respond to queries related to vehicle sales. However, identifying an intent category associated to "purchasing a vehicle over the Internet", may be very useful when responding to a user query.

Classifying queries according to their intent category changes the relationship between unique queries 12 and their frequency 14. This is represented by curve 22 where a significantly larger portion 20 of all received queries can be classified by a relatively small number of intent categories. Many of the outlier queries 18 previously located underneath curve 16 can be identified as having the same meaning or "intent" as some of the more frequently asked queries 20. Identifying the intent category for queries allow the search engine to provide more relevant responses to a larger percentage of queries while at the same time requiring substantially fewer resources to maintain the search engine. In other words, fewer responses can be used to adequately respond to a larger number of queries.

For example, a large number of queries received by a financial services enterprise may be related to 401K retirement plans. The queries related to 401Ks may be expressed by users in many different ways. For instance, "what is the current value of my 401K", "how much money is in my company retirement account", "show me the status of my 401K investments", etc. The information sought for each of these queries can be classified by the same intent category, namely: Intent Category=value of 401K. By classifying queries into intent categories, fewer associated responses have to be maintained.

FIG. 2 shows a computer network system 30 that includes an enterprise 32 that has one or more enterprise servers 34 and one or more enterprise databases 36. As described above, the enterprise 32 may be an online retailer that sells books and other retail items. In this example, the enterprise database 36

may contain price lists and other information for all of the books and other merchandise available for purchase. In another example, the enterprise 32 may be associated with a car dealership or financial institution and the enterprise database 36 could include vehicle or financial information, respectively. These are, of course, just examples, and any type of business or entity can be represented as enterprise 32.

Other web servers 26 may operate outside of the enterprise 32 and may include associated files or other web content 28. Examples of content stored in enterprise database 36 and in web server 26 may include HyperText Markup Language (HTML) web pages, Portable Document Format (PDF) files, Word® documents, structured database information or any other type of electronic content that can contain essentially any type of information.

Information in database 36 may be stored in a structured preconfigured format specified for the enterprise 32. For example, a book or vehicle price list may be considered structured content. The enterprise 32 may also generate and store specific intent responses 49 either in enterprise database 36 or on enterprise server 34 that are associated with specific intent categories 50. Other information that is contained in enterprise database 36, or contained on other web servers 26, may be considered non-structured content. This may include HTML web pages, text documents, or any other type of free flowing text or data that is not organized in a preconfigured data format.

A query 46 (e.g., electronic text question) may be initiated by a user from a terminal 25 through a User Interface (UI) 40. The terminal 25 in one example may be a Personal Computer (PC), laptop computer, wireless Personal Digital Assistant (PDA), cellular telephone, or any other wired or wireless device that can access and display content over a packet switched network. In this example, the query 46 is initiated from the UI 40 and transported over the Internet 48 to the enterprise server 34. For example, query 46 may be a question sent to a bank asking: Query="what is the current interest rates for CDs".

The enterprise server 34 operates a novel intent based search engine 35 alternatively referred to as an automated response system. The search engine 35 provides electronic responses, answers, and/or content pursuant to electronically submitted queries 46. The intent based search engine 35 uses a set of predetermined intent categories 50, one or more ontologies 52, and an Intelligent Matching Language (IML) engine 53 to identify the intent category 51 for query 46 and then provide an associated intent response 49.

The intent analysis is described in more detail below and converts the relatively flat query vs. frequency relationship curve 16 previously shown in FIG. 1 into the steeper query intent vs. frequency relationship curve 22. This results in the intent based search engine 35 presenting a more relevant intent based response 44 for electronically submitted question 42 while at the same time requiring a relatively low number of intent responses 49 for responding to a large number of unique queries 46. Accordingly, fewer resources have to be maintained by the intent based search engine 35.

The search engine 35 receives queries 46 from the UI 40 resulting from a question 42 entered by a user. The search engine 35 attempts to match the meaning or "intent" of the query 46 with preconfigured intent categories 50 using an intelligent matching language engine 53 and ontologies 52. The intent based search engine 35 then identifies one of the intent based responses 49 associated with the identified query intent category 51. The intent responses 49 may be preconfigured content or network links to information responsive to associated intent categories 50. The intent responses 49 can

5

also include any structured and/or non-structured content in the enterprise database 36 or on different web servers 26 that the intent based search engine 35 associates with the identified intent category 51. The identified information is then sent back to the UI 40 as intent based response 44.

The enterprise server 34 can include one or more processors that are configured to operate the intent based search engine 35. The operations performed by the intent based search engine 35 could be provided by software computer instructions that are stored in a computer readable medium, such as memory on server 34. The instructions are then executed by one or more of the processors in enterprise server 34. It should also be understood that the examples presented below are used for illustrative purposes only and the scope of the invention is not limited to any of the specific examples described below.

In one embodiment, the intent categories 50 are represented using a natural language, such as used in the IML engine 53. Using a natural language allows a system administrator to more easily create, delete, and modify intent categories 50. For example, the administrator can more easily identify which characteristics in an intent category need to be changed to more effectively classify a group of queries with a particular intent category. This is more intuitive than present information retrieval systems that use statistical analysis to classify queries into different categories.

Referring to FIGS. 2 and 3A, some of the different operations are described that may be performed by the intent based search engine 35. The search engine 35 may identify "a priori", the most frequently queried intent categories 50 (FIG. 2), and automatically display associated intent responses on the enterprise website. For example, 5-10% of the queries received by a financial service enterprise may contain questions related to retirement accounts. Accordingly, static links to web pages containing retirement account information may be presented on the home webpage for the financial institution prior to the user ever entering a question. This is referred to as pre-query based responses 60.

The intent based search engine 35 provides other types of responses according to the type of information that can be derived from the received queries. For example, there may be a set of around 100 intent categories 50 and corresponding intent responses 49 that address 60-70% of all unique queries received by a particular enterprise 32 (FIG. 2). A set of intent categories 50 are either manually or automatically derived based on previously received query information that cover this large percentage of unique queries. A set of intent responses 60 or 62 are then created that respond to the most frequently queried intent categories 50. The search engine 35 then attempts to match received queries with one of these frequent intent categories 50 and, if successful, sends back the corresponding intent based responses 60 or 62 (FIG. 3A).

Any identified intent categories 50 can also be used to improve the relevance of any other information provided in response to the query. For example, the identified intent category 50 in FIG. 2 may be used to identify both preconfigured intent based responses 49 and/or used for conducting an additional document search for other content in enterprise database 36 (FIG. 2) or other web content 28 in other web servers 26. The identified intent category 50 can also be used to extend or limit the scope of a document search or used to change the rankings for documents received back from the search.

The search engine 35 may use the IML engine 53 and ontologies 52 to discover concepts and other enterprise specific information contained in the queries 64. Any concepts discovered during query analysis can be used to discover the

6

intent categories and associated intent based responses 62. However, the search engine 35 may not be able to identify an intent category 50 for some percentage of less frequently received queries. If no intent categories can be identified, the search engine 35 can still use any identified concepts to provide ontology based responses 64.

To explain further, the ontologies 52 shown in FIG. 2 may associate different words such as IRA, 401 K, Roth, retirement, etc., with the broader concept of retirement accounts. The Intelligent Matching Language (IML) engine 53 is used in combination with the ontologies 52 to identify and associate these different phrases, words, and word forms, such as nouns, adjectives, verbs, singular, plural, etc., in the queries with different concepts.

For example, the IML engine 53 may receive a question asking about the price of a book but that does not necessarily include the symbol "\$", or use the word "dollar". The IML engine 53 may use ontologies 52 to associate the symbol "\$" and the word "dollar" with the words "Euro", "bucks", "cost", "price", "Yen", etc. The IML engine 53 then applies concepts such as <dollar> or <price> to the query 46 to then identify any words in the query associated with the <dollar> or <price> concepts. The identified concepts, words, etc., identified using IML engine 53 and ontology 52 are then used by the intent based search engine 35 to search for a relevant response.

One example operation of an IML engine 53 is described in patent application Ser. No. 10/820,341, filed Apr. 7, 2004, entitled: AN IMPROVED ONTOLOGY FOR USE WITH A SYSTEM, METHOD, AND COMPUTER READABLE MEDIUM FOR RETRIEVING INFORMATION AND RESPONSE TO A QUERY, which is herein incorporated by reference.

The intent based search engine 35 may also combine conventional keyword analysis with other intent and ontology based query analysis. If the search engine 35 does not identify any intent categories or ontology based concepts with the query, keyword based responses 66 may be provided based solely on keyword matching. For example, any content containing the same keywords used in the query 46 can be provided to the UI 40.

Thus, when no intent category can be determined, the search engine 35 may still use the domain based knowledge from the ontologies 52 to discover the most relevant responses 64. Alternatively, when the domain based knowledge does not provide any further understanding as to the meaning of the query, keyword analysis is used to provide keyword based responses 66. This is, of course, just one example of different combinations of intent, ontology concepts, and keyword analysis that can be performed on a query to provide different intent based responses 62, ontology based responses 64, and keyword based responses 66.

The intent based search engine 35 may conduct all of this intent, ontology and keyword analysis at the same time and then provide responses based on the highest level of query understanding. The intent based search engine 35 can use any other type of word, phrase, sentence, or other linguistic analysis, to determine the intent, concepts and words in the query. Similarly, any number of intent categories 50 and intent responses 49 may be used by the intent based search engine 35 and may cover any percentage of the unique queries 60, 62, 64, and 66 received by the enterprise server 34.

As described in more detail below, the intent based search engine 35 allows more efficient administration of an enterprise information system. For example, the most frequently referred to intent categories can be identified and associated intent responses derived. This provides a substantial advan-

tage over existing search engine administration where little or no ability exists for classifying multiple different queries with the same associated response. Similarly, the administrator can more efficiently add, delete, update, and/or modify the most relevant intent categories. In other words, the administrator is less likely to waste time generating or maintaining responses for infrequently received or irrelevant queries. This again is described in more detail below.

FIG. 3B shows in more detail how the search engine 35 identifies an intent category 51 for a query 46 and then provides an associated intent response 44. The query 46 received by the intent based search engine 35 is first analyzed by IML engine 53. The IML engine 53 uses natural language linguistic analysis to match the query 46 with one of the intent categories 50. One or more ontologies 52 are used that associated different words, phrases, etc., with different concepts that may be industry specific for the enterprise. For example, login, permission, password, passcode, etc., may all be associated with an <account information> concept that is regularly referred to by users accessing the enterprise network.

The IML 53 uses the ontologies 52, as well as other natural language associations, when trying to match the query 46 with one of the intent categories 50. When the query 46 is matched with one of the intent categories 50, the search engine 35 identifies the intent response 44 associated with the identified intent category 51. The identified intent response 44 is then displayed to the user.

For example, the following queries may either be received by the intent based search engine 35 in FIG. 2.

Queries	Intent Category
How do I change my password?	Change Password
I want to change my password	
How do I update my pass word?	
Is there a form for modifying passcodes?	
Change my password	
Need to change my secret code	
Is there a password change form?	

As seen above, each of these queries is associated with the same "change password" intent category 51. The search engine 35 may match each of these questions with the same intent category 51 and provide a corresponding intent response 44.

Intent Management Tool

FIG. 4 shows an intent management tool 67 that can be used to identify the most frequently queried intent categories 69A, identify the least frequently queried intent categories 69B, generate new intent categories 69C, identify queries 69D that do not match any existing intent categories, generate intent category hierarchies 69E, and/or assign and identify parameters to intent categories 69F.

The intent management tool 67 receives queries 68 that have been logged for some period of time by the enterprise server 34 (FIG. 2). The intent management tool 67 then uses existing intent categories 50, the intelligent matching language engine 53, and ontologies 52 to identify different information related to the logged queries 68 and the intent categories 50. For example, it may be beneficial to a website administrator to know which intent categories 69A match the most logged queries 68 or which intent categories 69B match the fewest logged queries 68. The intent management tool 67 can also automatically generate a new intent category 69C or identify queries 69D that do not match any existing intent categories. An intent category hierarchy 69E can be created

that is used for providing alternative responses to received queries. Multiple different parameters 69F can also be identified and assigned to different intent categories and then used for associating the intent categories with different intent responses.

FIG. 5A shows one example of how the intent management tool 67 generates a new intent category 79, or how the intent based search engine 35 matches a query with an existing intent category 79. Multiple different queries 70 may be received or logged by the enterprise server 34 (FIG. 2). Some of these queries 70A-70E may be associated with a same existing or non-existing intent category: CONTRACT EXPIRATION SUPPORT. A first one of the queries 70A however contains the question: "when does, my service expire". In a first spelling analysis stage 72, a natural language engine 71 used either by the intent based search engine 35 (FIG. 2) or intent management tool 67 (FIG. 4) checks the spelling in query 70A. The term "does" is not found in any of the ontologies 52 (FIG. 2). Accordingly, "does" is replaced with the closest match "do".

In a next punctuation and capitalization stage 73, punctuation is analyzed and a comma is removed that does not make sense. A speech analysis stage 74 analyzes nouns, verbs, etc., in the corrected question to generate an annotated question. For example, the word "do" is identified as a verb and the word "service" is identified as a noun phrase. Other words in an identified noun phrase may be removed. In a stem analysis stage 75, the search engine 35 or management tool 67 may add other forms of the same identified words to the annotated question. For example, the identified noun "service" could also be associated with "servicing", "serviced", etc.

In a concept analysis stage 76, the management tool uses the ontologies 52 to identify any concepts associated with the annotated question. For example, the word "service" may be associated with the concept <phone service contract> and the word "expire" may associated with the concept <end>. A linguistic analysis stage 77 then adds linguistic analysis results. For example, the annotated question may be determined to be related to an account, and accordingly restricted to related account and marketing information. In an intent analysis stage 78, an existing intent category is identified that matches the annotated question or a new intent category 79 is created for the previously annotated question. Similar linguistic analysis of questions 70B-70E may result in identifying the same existing intent category 79 or may be used along with query 70A to create a new intent category 79.

FIG. 5B describes in more detail how the intent categories are identified and created. A query 70F may ask the question: "I'm having trouble with my cell phone". The natural language engine 71 in combination one or more ontologies 52 are then used to conduct the concept analysis 76 and linguistic analysis 77 previously described in FIG. 5A. Different ontologies 52A-52C can be associated with different concepts. For example, ontology 52A is associated with the concept <trouble>, the ontology 52B is associated with pronouns and adjectives related to the concept <my>, and ontology 52C is associated with nouns and noun phrases related to the concept <cell phone>.

The natural language engine 71 uses ontologies 52 to identify different concepts 81A, 81B and 81C associated with the query 70F. The natural language engine 71 may identify a <my> concept 81A and a <trouble> concept 81C in query 70F. The natural language engine 71 may also identify the first <my> concept 81A as preceding a noun phrase 81B and also being located in the same sentence as the <trouble>

concept **81C**. This combination of concepts and associated sentence structure may be associated with a support query category **81E**.

The support query category **81E** may be associated with multiple different types of support intent categories and could even be identified as a parent intent category for multiple narrower support query intent categories in an intent category hierarchy. The natural language engine **71** uses the identified support query category **81E** along with the <cell phone> concept **81D** identified for the noun phrase **81B** to identify a cell phone support query intent category **81F** for query **70F**.

One system that conducts the linguistic analysis described in FIGS. **5A** and **5B** is the Inquire Matching Language described in patent application Ser. No. 10/820,341, filed Apr. 7, 2004, entitled: AN IMPROVED ONTOLOGY FOR USE WITH A SYSTEM, METHOD, AND COMPUTER READABLE MEDIUM FOR RETRIEVING INFORMATION AND RESPONSE TO A QUERY, which has already been incorporated by reference in its entirety. Of course, other types of natural language systems could also be used.

FIG. **6** explains further how the intent management tool **67** in FIG. **4** can be used to update intent categories. In operation **80**, previous queries are logged for some time period. For example, all of the queries for the past week. In operation **82**, the intent management tool compares the logged queries with existing intent categories. Any intent categories matching more than a first threshold number of logged queries may be identified in operation **83A**. Matching logged queries with existing intent categories can be performed in a similar manner as described above in FIGS. **5A** and **5B**. The intent responses for any identified intent categories in operation **83A** may then be posted on an enterprise webpage in operation **83B**. For example, 20% of the logged queries may have been associated with "contract expiration support" questions. If the threshold for adding an intent response to the enterprise web page is 15%, then a link to information relating to "contract expiration support" may be posted on the enterprise home web page in operation **83B**.

Optionally, the intent management tool may in operation **84A** identify information currently displayed or listed on the enterprise webpage that have an associated intent category that does not match a second threshold number of logged queries. Information associated with intent categories below this second threshold may be removed from the enterprise webpage in operation **84B**. For example, the enterprise home web page may currently display a link to an interest free checking promotion. If the number of logged queries matching an "interest free checking" intent category are below the second lower threshold, such below 1% of all logged queries, the "interest free checking" link or information can be identified by the intent management tool **67** and then either manually or automatically removed from the enterprise web page.

This provides a valuable system for promoting different services or products to users. For example, as described above, the intent management tool **67** can be used to determine that the "interest free checking" promotion is of little interest to customers. Alternatively, the same intent management tool **67** can determine that a "home refinance" promotion associated with a "home refinance" intent category has a substantially larger number of matching queries. Accordingly, a website administrator can quickly replace the interest free checking promotion with the home refinance promotion on the enterprise home web page.

In operation **85A**, the software executing the intent management tool **67** may automatically identify frequently queried intent categories that have no associated intent response. For example, the intent management tool **67** may identify

intent categories with no currently associated intent response that match a third threshold number of logged queries. In operation **85B**, the intent management tool **67** asks the administrator to identify an intent response for any identified intent categories. The intent responses can be information such as a web page and/or links to information on a web page that is responsive to the identified intent category. The intent responses input by the administrator are then assigned to the associated intent categories by the intent management tool **67** in operation **85C**.

In yet another operation **86A**, the software operating the intent management tool **67** may identify related queries with no associated intent categories. For example, a group of queries may be identified that are all related with a same financial service promotion but that currently have no assigned intent category. The largest number of related queries with no associated intent category may be identified first, and then lower numbers of related queries listed, etc. Alternatively, the intent management tool **67** can be configured to only list related queries over some predetermined threshold number.

The intent management tool in operation **86B** asks the administrator to identify an intent category for the group of identified queries. Alternatively, the common information identified in the group of queries may be used as the intent category. In operation **86C**, the intent management tool **67** then asks the user to identify an intent response for the identified intent category.

FIG. **7** shows one way the intent management tool **67** can be used to update existing intent categories. In operation **90**, queries are logged for some period of time in the same manner described above in FIG. **6**. In operation **92**, the intent management tool **67** identifies logged queries that do not match any current intent categories. One or more new intent categories are then created for the non-matching queries in operation **94**. The new intent categories are either manually generated by the administrator or automatically generated by a natural language engine **71** as described above in FIGS. **5A** and **5B**.

The new intent categories are then run against the logged queries in operation **96**. This operation may be iterative. For example, the number of matching queries is identified in operation **98**. If the number of queries matching the new intent category is below some threshold, such as a predetermined percentage of the logged queries, the new intent category may be refined in operation **97** and then compared again with the logged queries in operation **96**. For example, the administrator may modify certain words in the intent category that may cause more matches with the logged queries. When the new intent category matches more than some threshold number of logged queries in operation **98**, IML is created in operation **99** that automatically matches the same logged queries with the new intent category.

In another embodiment, the new intent category may also be applied to other query logs to verify accuracy. For example, a query log from another time period may be applied to the newly created intent category. The operation described above for generating new intent categories can also be used when generating the initial intent categories for an enterprise.

It should also be understood that an industry expert may be used to review the logged queries and then manually generate useful intent categories based on the results from the intent management tool **67**. The fact that the intent categories are "useful" is worth noting. Some clustering algorithms may generate information categories that, for example, may be too broad to really provide useful information. For example, as described above, a clustering algorithm may identify queries

all related to “email”. However, providing and supporting a general email intent category may be of little relevance when trying to provide responses to queries directed to an online financial institution.

The industry expert can first derive pertinent intent categories and then refine the derived intent categories to optimize the number of queries matches. This ensures that intent categories are useful and are relevant to the actual queries submitted by users. The optimized intent categories are then used by the search engine to identify query meaning. Alternatively, all or part of the intent category generation can be automated using the intent discovery tool 67 as described above in FIGS. 5A and 5B and as described in further detail below.

The intent discovery tool 67 also allows the web site administrator to identify queries that do not correspond with current intent categories. For example, users may start using new terminology in queries referring to a new service or product. The intent discovery tool 67 can identify these queries that do not match existing intent categories and then either modify an existing related intent category or create a new intent category that matches the identified queries.

Intent Hierarchy

FIG. 8 shows how hierarchies can be associated with intent categories. In this example, a group of queries 100 are all associated with a retirement plan research intent category 110. Either manually or through the intent management tool 67 in FIG. 4, a Roth intent category 102 is derived for a first group of queries 100A, a Regular IRA intent category 104 is created for a second group of queries 100B, and a 401K intent category is derived for a third set of queries 100C.

Again either manually by an industry expert, or automatically with the management tool 67, an intent hierarchy 126 is derived for the intent categories 102-110. For example, a parent “IRA” intent category 108 is derived for intent categories 102 and 104. In addition, a parent “Retirement Plan Research” intent category 110 is derived for intent categories 108 and 106.

This intent hierarchy 126 can be derived in a variety of different ways, but in one example could use clustering analysis as described in more detail below in FIG. 15. A hierarchy tag can then be assigned to the intent categories to control what responses are automatically presented to a user.

To explain further, FIG. 9 shows how the intent hierarchy 126 is used in combination with an identified intent category 102. The intent based search engine may receive a query 120 that matches the Roth intent category 102 previously described in FIG. 8. Accordingly, the search engine displays an intent response 128A associated with the identified “Roth” intent category 102.

However, the intent category 102 can also be assigned a tag 124 that directs the search engine to display responses for any parents of the “Roth” intent category 102. Accordingly, by selecting tag 124, the search engine refers to the intent category hierarchy 126 to identify any parents of “Roth” intent category 102. In this example, the “IRA” intent category 108 and the “Retirement Plan Research” intent category 110 are identified as parents. Accordingly, intent responses 128B and 128C associated with intent categories 108 and 110, respectively, are also displayed in response to query 120.

Notice that in this example, the intent response 128C associated with parent intent category 110 includes a promotional advertisement for opening a 401K account. Since the “Roth” intent category 102 and the “401K” intent category 106 both have a common parent 110, the enterprise can use tag 124 to promote services, products, or present other information to

users that is related to a common broader subject matter than what is actually contained in query 120.

This intent hierarchy feature provides a powerful tool for providing relevant information responsive to a query. For example, a user may not necessarily know they are seeking information related to a 401K account. However, the user is aware of IRA accounts. Intent hierarchy tag 124 allows the search engine to automatically associate a question related to IRAs with an intent category related to 401K accounts based on the classification of both IRA and 401K accounts under the same “Retirement Plan research” parent intent category 110. Thus, the user may receive some relevant 401K information under the “Retirement Plan research” intent category 110 without ever using the word “401K” in the query 120. This also has promotional and advertising advantages. For example, the enterprise can notify any user sending any type of retirement plan related query of a new 401K promotion.

The intent hierarchy tag 124 can consist of a pointer to an associated intent hierarchy 126 as shown in FIG. 9. The intent hierarchy tag 124 can also be used to direct the search engine to display intent responses associated with child intent categories or associated with other intent categories not contained in the same hierarchy.

Parameterized Intent Categories

FIG. 10 shows another embodiment of the intent based search engine 35 that allows an administrator or user to associate different parameters with intent categories. The intent management tool 67, for example, may process logged queries 68. In this example, the intent management tool 67 either identifies or creates a “vehicle research” intent category 130 and may then assign different intent responses 134 to the intent category 130 using parameters 135.

The management tool 67 automatically compares the intent category 130 with one or more ontologies 133 and determines that the word “vehicle” 131 in the intent category 130 is associated with the <vehicle> concept 132A in ontology 133. The management tool 67 may then present the user with a drop down menu or, some other type of display, that shows the different concepts or other words or phrases associated with the <vehicle> concept 132A in ontology 133. In this example, the concepts 132A-132E in ontology 133 are displayed as parameters 137A-137E, respectively. The parameters 137A-137E may include pointers to associated intent responses 134A-134E, respectively.

The administrator can select which of the parameters 137A-137E (pointers to intent responses 134) to associate with the intent category 130. In this example, the administrator at least selects minivan parameter 137B. The search engine 35 will then use the assigned parameter 137B to provide additional responses to an associated query. For example, the search engine 35 may later receive a query 139 associated with the vehicle research intent category 130. The search engine 35 identifies the selected parameter 137B assigned to intent category 130 and accordingly displays the intent responses 134B.

In another embodiment, the intent parameters 135 may also cause the search engine to display responses for any associated parent concepts. For example, a query may be associated with a minivan research intent category. A parameter 135 can be assigned to the minivan intent category that causes the search engine to provide responses associated with any of the broader concepts in ontology 133, such as a response related to the vehicle research intent category 130.

The selection of different parameters 135 can similarly be performed by a user. For example, the search engine 35 may initially display the intent category 130 to the user based on a

13

received query. The user can then be provided with the same list of different parameters **137A-137E** associated with the ontology **133**. The user then selects which intent responses **134** to display by selecting any combination of parameters **137A-137D**.

It is worth noting that the intent category hierarchy described above in FIGS. **8** and **9** and the intent parameters shown in FIG. **10** may be useful in classifying different types of queries. For example, the intent hierarchies in FIGS. **8** and **9** may be better at classifying queries that include more verbs, and the intent parameters in FIG. **10** may be better at classifying queries that include more nouns. For example, questions related to specific types of products may include more nouns while questions related to services or user activities may include more verbs. Of course, these are just examples and either the intent category hierarchy or the intent parameters can be used for any type of query.

Generating New Intent Parameters

Referring still to FIG. **10**, the intent management tool **67** can also be used for identifying new intent parameters **140**. The intent management tool **67** may identify a large group of queries all matching intent category **130** but that do not match any of the existing parameters **135** or associated concepts **132** in ontology **133**. For example, a group of queries may all be associated with a new minivan model C that is not currently identified in ontology **133**.

The intent management tool **67** suggests adding a new parameter **137F** to parameter list **135** that is associated with the identified minivan model C. Upon selection, parameter **137F** is added to parameter list **135**. The intent management tool **67** may also ask the administrator to add any other synonyms associated with the new model C parameter **137F** and provide an associated intent response **134F**. In addition, the intent management tool **67** may update ontology **133** to include a new model C concept **132F** underneath the minivan concept **132B**.

User Classification

The intent management tool **67** can also assign different “user” related parameters to intent categories. This allows the intent based search engine to associate particular intent responses or search engine actions with different types of users. For example, it may be desirable to automatically initiate a phone call to any long term user that has sent a query associated with terminating an existing account. In another scenario, it may be desirable for the search engine to track the number of times particular users send queries associated with particular intent categories. The search engine can then send intent responses based on the tracked frequency.

Referring to FIG. **11**, any of these different user associated parameters are assigned to particular intent categories by the administrator using the intent management tool **67**. The intent based search engine **35** may then receive a query in operation **150**. The search engine identifies an intent category for the query in operation **152** and identifies any user parameters that may be associated with the identified intent category in operation **154**.

The search engine in operation **156** conducts any user operation according to the identified user parameters. For example, the user parameter may direct the search engine in operation **158** to track the user query frequency and then classify the user according to the identified frequency. This could be used for providing special promotional materials to high frequency users. Accordingly, the user parameter may direct the search engine in operation **159** to display certain intent responses to the user according to the user classification. The user classifications can also be based on factors

14

unrelated to user query frequency. For example, the user classifications may be based on how long the user has been signed up on the enterprise website; priority membership status, such as a platinum membership, geographic region, age, or any other user demographic.

Intent Discovery

Clustering algorithms are used for statistically associating together different information. A set of features are input into the clustering algorithm which then groups together different information according to the features. These types of conventional clustering algorithms are known to those skilled in the art, and are accordingly not described in further detail.

The present intent discovery scheme may provide different concepts to the clustering algorithm as features that then allow the clustering algorithm to more effectively cluster together related queries. The features provided to the clustering algorithm can be any combination of words, stems, tokens, phrases, concepts, intent categories, etc. FIG. **13** describes this intent discovery scheme in more detail. A set of queries **175** may be input to a clustering engine **186**. As opposed to conventional keyword clustering, the clustering engine **186** is given at least a partial set of features **184** associated with the concepts in an enterprise specific ontology **183**. For example, the stems, tokens, phrases, and/or concepts in ontology **183** may all be associated with the concept “family vehicle”.

The clustering engine **186** analyzes the queries **175** with respect to the ontology based features **184**. Some higher order concepts, such as the concept “family vehicle” may get a larger weight when the queries **175** are clustered than lower order concepts, such as “vehicle models”. The clustering engine **186** outputs names **188** for the identified clusters that, for example, may comprise a string of the most common terms and highest order concepts in the clusters.

Then either through a manual or automated process, IML expressions **190** are created that match the queries in the identified clusters with a particular intent category. The intent categories may use some or all of the terms from the cluster names. For example, the string of most common terms **192** contained in queries **182** may be used in the IML expression **190** to identify station wagon queries **182**. Other concepts in ontology **183** can also be used in the IML expression **192** to help classify the station wagon queries **182**.

Referring to FIG. **13**, the above clustering scheme can also be used to further improve or automate intent classification. For example, the intent management tool **67** described in FIG. **4** may be used in operation **190** to identify any of the logged queries that do not match any of the existing intent categories. In operation **192**, the identified queries are submitted to the clustering engine **186** in FIG. **12**. In operation **194**, features from one or more of the ontologies **183** in FIG. **12** are also fed into the clustering engine **186**. The intent management tool **67** receives the names identified by the clustering engine in operation **196** and uses the cluster names and the identified clustered queries to generate new intent categories in operation **198**.

Referring to FIG. **14**, the intent discovery scheme can also be used to create intent hierarchies. For example, intent category **200** for “family vehicles” and intent subcategory **201** for “minivans” have already been created. However, the intent discovery scheme described above may have discovered three new intent categories **202A-202C**.

The intent management tool **67** may compare the queries matching multiple intent categories to determine where the new intent categories **202A-202C** should be located in the hierarchy. For example, the intent management tool **67** may

15

discover that all of the queries matching new intent category **202C** are a subset of the queries matching existing parent intent category **200**. Further, the intent management tool **67** may also determine that the queries matching new intent category **202C** do not, or rarely, overlap with the queries matching “minivan” intent category **201**. Accordingly, the intent management tool **67** locates new intent category **202C** as a direct child of intent category **200**.

It may be determined that the queries matching the other new intent categories **202A** and **202B** are a subset of the queries matching existing intent category **201**. Accordingly, new intent categories **202A** and **202B** are assigned as direct descendants of intent category **201**. The intent management tool **67** may also identify new parameters for an existing intent category as described above in FIG. **10**.

Response Parameters

FIG. **15** shows another type of parameter that can be assigned to different intent responses. An intent response **220** may comprise a template associated with a particular general category of questions. For example, the intent response **220** may be associated with an intent category related to buying and viewing a vehicle. Instead of creating a separate intent response for every specific model of vehicle that a user may ask about, the intent response **220** may include parameters **222** and **224** that are associated with specific information elements within the query.

For example, response parameter **222A** may be associated with price information **228A** for a particular minivan model and response parameter **222B** may be associated with price information **228C** for a particular station wagon model. Similarly, response parameter **224A** may be associated with image information **228B** for the minivan and response parameter **224B** may be associated with image information **228D** for the station wagon.

The intent based search engine **35** receives the query **230** and conducts the linguistic analysis described above to determine an associated intent category. The identified intent category is associated with intent response **220**. The search engine **35** then compares elements in the query **230** with the response parameters **222** and **224** to determine what additional response elements **228** to insert into intent response **220**.

In this example, the search engine matches the <minivan> concept parameters **222A** and **224A** in intent response **220** with the word minivan in query **230**. Accordingly, the response elements **228A** and **228B** in table **226** are displayed with the intent response **220** on user interface **232**. The response parameters allow an almost identical intent response **220** to be generated for all of the queries within a particular intent category and then automatically customize the intent response **220** for different query elements.

The system described above can use dedicated processor systems, micro controllers, programmable logic devices, or microprocessors that perform some or all of the operations. Some of the operations described above may be implemented in software and other operations may be implemented in hardware.

For the sake of convenience, the operations are described as various interconnected functional blocks or distinct software modules. This is not necessary, however, and there may be cases where these functional blocks or modules are equivalently aggregated into a single logic device, program or operation with unclear boundaries. In any event, the functional blocks and software modules or features of the flexible interface can be implemented by themselves, or in combination with other operations in either hardware or software.

16

Having described and illustrated the principles of the invention in a preferred embodiment thereof, it should be apparent that the invention may be modified in arrangement and detail without departing from such principles. I/We claim all modifications and variation coming within the spirit and scope of the following claims.

The invention claimed is:

1. A method for using a computer for identifying query intent, comprising:

receiving multiple different queries at an input of the computer;

performing a natural language analysis with the computer to identify the multiple different queries that use different natural language formations of words, phrases, and concepts to request a similar category of information;

identifying with the computer a same intent category for all of the identified queries requesting the similar category of information;

creating with the computer a same single natural language word or phrase for the identified same intent category;

using the same single natural language word or phrase to summarize in a same single natural language word formation a common generalized information request intent that describes a particular category of information requested by all of the multiple different queries, wherein the multiple different queries use multiple different natural language formations to request different types of information all summarized by the natural language word or phrase created for the same intent category;

associating an intent response formulated of natural language words or phrases with the identified intent category, wherein the intent response provides a single common applicable response to the particular category of information and to the different types of information requested by all of the multiple different queries in the same intent category;

receiving a new query;

comparing, with the computer, concepts and words in the natural language formations in the new query with concepts and words in the same single natural language word formation of the intent category;

displaying the intent response for the intent category when the concepts or words in the new query match the concepts or words in the natural language word formation of the intent category;

providing one or more ontologies that associate different concepts related to a particular industry, the ontologies configured to link a plurality of concepts at multiple different concept domain layers, wherein the concepts associated with the different domain layers are represented by natural language words and the concepts for higher concept domain layers are represented by natural language words having more specific terms than the natural language words representing the concepts associated with the lower domain layers;

identifying the concepts in the ontologies corresponding with a word or phrase in the queries;

using the identified concepts to identify the intent category for the queries;

identifying a new intent category for the new query by comparing a natural language phrase representing the new intent category with a natural language phrase in the new query;

identifying an intent hierarchy of multiple hierarchical intent categories associated with the identified new intent category wherein natural language phrases repre-

17

senting the hierarchical intent categories at higher parent domain layers of the intent hierarchy have more general terms than natural language phrases representing the hierarchical intent categories at lower child domain layers of the intent hierarchy, and wherein the intent hierarchy is independent and different from the hierarchy for the one or more ontologies;

displaying a new intent response corresponding with a location of the identified new intent category in the intent hierarchy;

determining if the new intent category has parent hierarchical intent categories located at a higher domain layer of the intent hierarchy; and

displaying intent responses corresponding to the parent hierarchical intent categories associated with the new intent category when a response tag is set in the new intent category.

2. The method according to claim 1 wherein the same single natural language phrase for the intent category does not exist in any of the multiple different queries and is not a natural language query for information.

3. The method according to claim 1 including:

analyzing with the computer how the identified concepts and other words in the multiple different queries are arranged in sentence structures to identify different natural language meanings; and

automatically generating with the computer the natural language word or phrase for the intent category, the natural language word or phrase both summarizing and paraphrasing the identified different natural language meanings for all of the multiple different queries into the same particular category of information.

4. The method according to claim 1 including:

logging the queries for a certain period of time;

identifying at least some most frequently queried intent categories for the logged queries; and

providing pre-query intent responses or links for at least a portion of the identified most frequently queried intent categories.

5. The method according to claim 4 including:

receiving a first group of queries that match the identified most frequently queried intent categories; and

providing pre-configured intent responses for the first group of received queries corresponding with the matching intent categories.

6. The method according to claim 5 including:

receiving a second group of queries that cannot be matched with the identified most frequently queried intent categories;

using a natural language engine and the one or more ontologies to identify concepts in the second group of non-matching queries;

comparing the identified concepts for the second group of queries with natural language words or phrases for other intent categories; and

providing predetermined responses to any of the second group of queries matching any of the other intent categories.

7. The method according to claim 6 including:

receiving a third group of queries that cannot be matched with any of the identified most frequently queried intent categories and cannot be matched with the concepts in the ontologies;

identifying keywords in the third group of queries; and

using the identified keywords to search for responses associated with the third group of queries.

18

8. The method according to claim 1 including:

receiving queries;

identifying intent categories for the received queries;

assigning parameters to at least some of the identified intent categories that correspond with the concepts in the ontologies; and

displaying intent responses corresponding with the assigned parameters for the identified intent categories.

9. The method according to claim 1 including:

automatically and dynamically creating with the computer a single natural language phrase of multiple different natural language words describing the identified intent category, wherein automatically and dynamically creating the phrase of the multiple different natural language words comprises:

automatically correcting spelling in the different queries;

automatically correcting capitalization and punctuation in the different queries;

automatically performing speech analysis in the queries to identify different types of words in the queries;

automatically performing stem analysis on the queries to add other forms of similar words to the queries;

automatically performing concept analysis on the queries using an ontology to identify different words that have similar meanings as the words in the queries;

automatically performing linguistic analysis on the queries to identify subject categories for the words in the queries; and

automatically performing intent analysis on the subject categories and concepts identified in the queries to automatically generate the natural language word or phrase summarizing the different natural language formations associated with the multiple different queries in the same intent category.

10. A search engine, comprising:

a processor configured to:

receive different queries and then conduct a linguistic analysis that identifies different concepts and linguistic characteristics for different natural language word formations in the different queries;

identify the received queries having similar information requests that use different natural language phrases to make the similar information requests, wherein the similar information requests are identified according to the identified concepts and linguistic characteristics for the different natural language word formations, and wherein the queries with similar information requests are classified under a same intent category;

extract natural language words or phrases or different combinations of natural language words or phrases from the different concepts and linguistic characteristics in the different natural language word formations in the different queries;

create with the processor a same single natural language word or phrase for the same intent category;

use the same single natural language word or phrase to summarize in a same natural language word formation a common generalized information request intent that describes a particular category of information requested by the different queries, wherein the different queries use the different natural language formations to request different types of information all summarized by the natural language word or phrase created for the same intent category; and

19

provide common information responses to the similar information requests classified under the same intent category; and

the processor further configured to:

provide a hierarchy of one or more ontologies that associate different concepts, the ontologies configured to link a plurality of concepts at multiple different concept domain layers, wherein the concepts associated with the different concept domain layers are represented by natural language words and the concepts for the different concept domain layers are represented by natural language words having different specificity of terms than the natural language words representing the concepts associated with other domain layers;

identify the concepts in the ontologies corresponding with the linguistic characteristics for the different natural language word formations in the received queries;

use the identified concepts to identify the intent categories for at least some of the received queries;

provide an intent hierarchy of the intent categories wherein natural language phrases representing the hierarchy of the intent categories at different domain layers of the intent hierarchy have different generalities of terms than natural language phrases representing the intent categories in other domain layers of the intent hierarchy;

identify different intent responses corresponding to locations of the intent categories in the intent hierarchy wherein the intent hierarchy is separate and different from the hierarchy for the ontologies used for identifying the concepts corresponding with the linguistic characteristics in the queries; and

display the intent responses corresponding with the locations of the intent categories identified in the intent hierarchy.

11. The search engine according to claim 10 wherein the processor is configured to operate an Intelligent Matching Language (IML) engine and use the one or more of the ontologies to identify the concepts and linguistic characteristics in the queries and compare the concepts and linguistic characteristics identified in the queries with concepts and linguistic formations in the natural language word formations or phrases representing different intent categories.

12. The search engine according to claim 11 including a memory that stores preconfigured intent responses for only a portion of most frequently queried intent categories and returns the preconfigured intent responses for the intent categories associated with the received queries.

13. The search engine associated with claim 12 wherein the processor is configured to conduct searches using concepts identified using the IML engine and the ontologies when no intent categories can be identified for the received queries and is further configured to conduct keyword searches when no concepts can be identified in the queries.

14. The search engine according to claim 10 wherein the intent category for the queries with similar information requests includes different parameters associated with the different intent responses.

15. The search engine according to claim 14 wherein the parameters are associated with the ontology concepts and the processor displays intent responses corresponding with the ontology concepts.

20

16. The search engine according to claim 14 wherein at least some of the parameters are associated with user frequencies indicating how frequently users use the search engine, wherein the processor is further configured to provide responses according to their identified user frequencies.

17. The search engine according to claim 14 wherein the different intent responses include response parameters associated with additional intent responses and the processor is configured to identify elements in the queries associated with the response parameters and display the associated additional intent responses.

18. An article comprising a device, having stored thereon instructions that, in response to execution by a processing device, result in:

receiving multiple different queries;

performing a natural language analysis to identify the different queries that use different natural language formations of words, phrases, and concepts to request a similar category of information;

identifying same intent categories for the queries requesting the similar categories of information;

creating a same single natural language word or phrase for the same intent categories;

using the same single natural language word or phrase to summarize, in a same natural language word formation, common generalized information requested by the different queries, wherein the queries use the different natural language formations to request different types of information summarized by the natural language word or phrase created for the same intent categories;

associating intent responses formulated of natural language words or phrases with the identified same intent category, wherein the intent responses provides a single common applicable responses to the particular category of information and to the different types of information requested by all of the multiple different queries in the identified same intent categories;

providing one or more ontologies configured to link a plurality of concepts at multiple different concept domain layers, wherein the concepts associated with the different domain layers are represented by natural language words and the concepts for different domain layers are represented by natural language words having different specificity of terms than the natural language words representing the concepts associated with lower concept domain layers;

identifying the concepts in the ontologies corresponding with the word or phrase in the queries;

using the concepts identified in the ontologies to identify the intent categories for the queries;

providing an intent hierarchy of multiple intent categories wherein natural language phrases representing the intent categories at different domain layers of the intent hierarchy have different generalities of terms than natural language phrases representing the intent categories in other domain layers of the intent hierarchy, and wherein the intent hierarchy is independent and different from the hierarchy for the one of more ontologies; and

displaying intent responses corresponding to the intent categories in the intent hierarchy associated with the queries.

* * * * *

UNITED STATES PATENT AND TRADEMARK OFFICE
CERTIFICATE OF CORRECTION

PATENT NO. : 7,747,601 B2
APPLICATION NO. : 11/464443
DATED : June 29, 2010
INVENTOR(S) : Cooper et al.

Page 1 of 1


It is certified that error appears in the above-identified patent and that said Letters Patent is hereby corrected as shown below:

At column 16, line 52, please replace "words" with --words,--.

At column 16, line 67, please replace "category" with --category,--.

At column 19, line 19, please replace "the intent" with --intent--.

Signed and Sealed this
Twenty-sixth Day of July, 2011

A handwritten signature in black ink, reading "David J. Kappos". The signature is written in a cursive, flowing style with a large initial "D" and a stylized "K".

David J. Kappos
Director of the United States Patent and Trademark Office