



US008082264B2

(12) **United States Patent**
Bierner

(10) **Patent No.:** **US 8,082,264 B2**
(45) **Date of Patent:** **Dec. 20, 2011**

(54) **AUTOMATED SCHEME FOR IDENTIFYING
USER INTENT IN REAL-TIME**

(75) Inventor: **Gann Alexander Bierner**, San Bruno,
CA (US)

(73) Assignee: **Inquira, Inc.**, Redwood Shores, CA
(US)

(*) Notice: Subject to any disclaimer, the term of this
patent is extended or adjusted under 35
U.S.C. 154(b) by 398 days.

5,386,556 A 1/1995 Hedin et al.
5,442,782 A 8/1995 Malatesta et al.
5,535,382 A 7/1996 Ogawa
5,608,624 A 3/1997 Luciw
5,625,814 A 4/1997 Luciw
5,694,523 A 12/1997 Wical

(Continued)

FOREIGN PATENT DOCUMENTS

WO WO 0188662 A2 11/2001
(Continued)

OTHER PUBLICATIONS

Zaiane et al., Finding Similar Queries to Satisfy Searches Based on
Query Traces. Advances in Object-Oriented Information Systems,
Lecture Notes in Computer Science. vol. 2426/2002. [online], 2002
[retrieved on Nov. 18, 2010]. Retrieved from the Internet:<
URL:http://www.springerlink.com/content/h2ukf5jflymp6dr7/
fulltext.pdf>. pp. 207-216.*

(Continued)

(21) Appl. No.: **11/959,307**

(22) Filed: **Dec. 18, 2007**

(65) **Prior Publication Data**

US 2008/0104037 A1 May 1, 2008

Related U.S. Application Data

(63) Continuation-in-part of application No. 10/820,341,
filed on Apr. 7, 2004, and a continuation-in-part of
application No. 11/464,443, filed on Aug. 14, 2006,
now Pat. No. 7,747,601.

(60) Provisional application No. 60/883,279, filed on Jan.
3, 2007.

(51) **Int. Cl.**
G06F 17/30 (2006.01)

(52) **U.S. Cl.** 707/767; 707/768; 707/771; 715/816

(58) **Field of Classification Search** 707/705-708,
707/771, 766-768; 715/816

See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

4,956,774 A 9/1990 Shibamiya et al.
5,321,833 A 6/1994 Chang et al.
5,357,436 A 10/1994 Chiu

Primary Examiner — John Cottingham

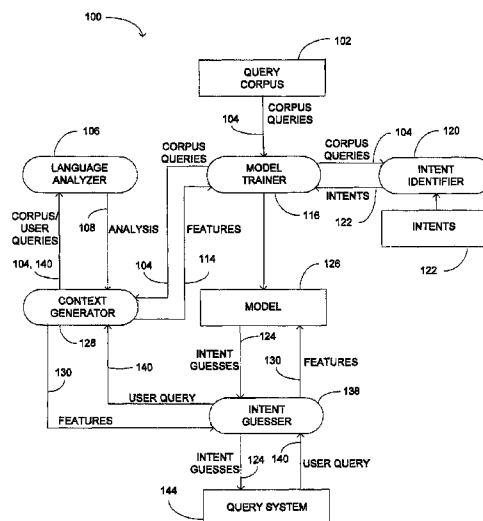
Assistant Examiner — James E Richardson

(74) *Attorney, Agent, or Firm* — Kilpatrick Townsend &
Stockton LLP

(57) **ABSTRACT**

An intent guessing system receives partial user queries as
they are entered by the user. The partial user queries are
compared with different intents derived from previously
logged queries. Guesses are made as to which of the intents
are associated with the partial user query. The intent guesses
are then provided as responses to the user query. Features are
identified for the earlier logged queries and associated with
the derived intents. The derived intents and associated fea-
tures are then used to identify intents for the partial user
queries.

15 Claims, 6 Drawing Sheets



U.S. PATENT DOCUMENTS

5,694,546 A 12/1997 Reisman
 5,742,816 A 4/1998 Barr et al.
 5,794,050 A 8/1998 Dahlgren et al.
 5,812,865 A 9/1998 Theimer et al.
 5,826,269 A 10/1998 Hussey
 5,848,399 A 12/1998 Burke
 5,873,056 A 2/1999 Liddy et al.
 5,873,076 A 2/1999 Barr et al.
 5,873,080 A 2/1999 Coden et al.
 5,878,423 A 3/1999 Anderson et al.
 5,884,302 A 3/1999 Ho
 5,890,152 A 3/1999 Rapaport et al.
 5,893,091 A 4/1999 Hunt et al.
 5,897,622 A 4/1999 Blinn et al.
 5,901,287 A 5/1999 Bull et al.
 5,913,215 A 6/1999 Rubinstein et al.
 5,933,822 A 8/1999 Braden-Harder et al.
 5,948,054 A 9/1999 Nielsen
 5,963,940 A 10/1999 Liddy et al.
 5,966,695 A 10/1999 Melchione et al.
 5,974,412 A 10/1999 Hazlehurst et al.
 5,987,454 A 11/1999 Hobbs
 5,995,921 A 11/1999 Richards et al.
 6,006,225 A 12/1999 Bowman et al.
 6,016,476 A 1/2000 Maes et al.
 6,021,403 A 2/2000 Horvitz
 6,026,388 A 2/2000 Liddy et al.
 6,028,601 A 2/2000 Machiraju et al.
 6,038,560 A 3/2000 Wical
 6,052,710 A 4/2000 Saliba et al.
 6,061,057 A 5/2000 Knowlton et al.
 6,070,149 A 5/2000 Tavor et al.
 6,078,914 A 6/2000 Redfern
 6,151,600 A 11/2000 Dedrick
 6,208,991 B1 3/2001 French et al.
 6,233,547 B1 5/2001 Denber
 6,260,035 B1 7/2001 Horvitz et al.
 6,278,996 B1 8/2001 Richardson et al.
 6,282,534 B1 8/2001 Vora
 6,285,998 B1 9/2001 Black et al.
 6,370,535 B1 4/2002 Shapiro
 6,385,592 B1 5/2002 Angles et al.
 6,393,479 B1 5/2002 Glommen et al.
 6,401,084 B1 6/2002 Ortega et al.
 6,401,094 B1 6/2002 Stemp et al.
 6,466,899 B1 10/2002 Yano
 6,471,521 B1 10/2002 Dornbush et al.
 6,477,551 B1 11/2002 Johnson et al.
 6,480,843 B2 11/2002 Li
 6,564,213 B1 * 5/2003 Ortega et al. 1/1
 6,584,464 B1 6/2003 Warthen
 6,675,159 B1 1/2004 Lin et al.
 6,766,320 B1 7/2004 Wang et al.
 6,907,414 B1 6/2005 Parnell
 6,928,425 B2 8/2005 Grefenstette et al.
 6,941,301 B2 9/2005 Ferguson et al.
 6,944,711 B2 9/2005 Mogi et al.
 6,957,213 B1 10/2005 Yuret
 7,024,400 B2 4/2006 Tokuda et al.
 7,035,864 B1 4/2006 Ferrari et al.
 7,035,910 B1 4/2006 Dutta et al.
 7,111,290 B1 9/2006 Yates
 7,177,795 B1 2/2007 Chen et al.
 7,181,731 B2 2/2007 Pace
 7,209,921 B2 4/2007 Pace
 7,254,806 B1 8/2007 Yates
 7,428,541 B2 9/2008 Houle
 7,668,850 B1 2/2010 Bierner
 7,672,951 B1 3/2010 Bierner
 7,676,517 B2 * 3/2010 Hurst-Hiller et al. 707/713
 2001/0014888 A1 8/2001 Tsuchida et al.
 2001/0047270 A1 11/2001 Gusick et al.
 2002/0049738 A1 4/2002 Epstein
 2002/0051020 A1 5/2002 Ferrari
 2002/0078090 A1 6/2002 Hwang et al.
 2002/0103809 A1 8/2002 Starzl et al.
 2003/0014403 A1 * 1/2003 Chandrasekar et al. 707/5
 2003/0018512 A1 1/2003 Dortmans

2003/0037073 A1 2/2003 Tokuda et al.
 2003/0101153 A1 5/2003 Francis et al.
 2003/0144994 A1 7/2003 Wen et al.
 2003/0204435 A1 10/2003 McQuilkin et al.
 2003/0217052 A1 11/2003 Rubenczyk et al.
 2003/0233224 A1 12/2003 Marchisio et al.
 2004/0024739 A1 2/2004 Copperman et al.
 2004/0167870 A1 8/2004 Wakefield et al.
 2005/0080775 A1 4/2005 Colledge et al.
 2005/0120045 A1 6/2005 Klawon
 2005/0149496 A1 * 7/2005 Mukherjee et al. 707/3
 2005/0198563 A1 9/2005 Kristjansson
 2005/0203878 A1 * 9/2005 Brill et al. 707/3
 2005/0234969 A1 10/2005 Mamou et al.
 2005/0267871 A1 12/2005 Marchisio et al.
 2006/0059073 A1 3/2006 Walzak
 2006/0074836 A1 4/2006 Gardner
 2006/0095326 A1 5/2006 Sandhu et al.
 2006/0106769 A1 * 5/2006 Gibbs 707/3
 2006/0122979 A1 6/2006 Kapur et al.
 2006/0136403 A1 6/2006 Koo
 2006/0179404 A1 8/2006 Yolleck et al.
 2006/0253427 A1 * 11/2006 Wu et al. 707/3
 2006/0259483 A1 11/2006 Ozana
 2006/0265232 A1 11/2006 Katariya et al.
 2007/0033116 A1 2/2007 Murray
 2007/0208738 A1 * 9/2007 Morgan 707/6
 2007/0282769 A1 12/2007 Bierner
 2008/0189163 A1 8/2008 Rosenberg
 2008/0215976 A1 9/2008 Bierner
 2008/0243761 A1 10/2008 Guo
 2008/0294651 A1 11/2008 Masuyama et al.
 2009/0077047 A1 3/2009 Cooper
 2009/0083224 A1 3/2009 Dettinger et al.
 2009/0089044 A1 4/2009 Cooper

FOREIGN PATENT DOCUMENTS

WO WO 0235376 A2 5/2002
 WO WO/2007/134128 11/2007
 WO WO/2008/022150 2/2008
 WO WO/2008/067316 5/2008

OTHER PUBLICATIONS

Baeza-Yates et al. "Query Recommendation Using Query Logs in Search Engines" Mar. 2004, Current Trends in Database Technology—EDBT 2004 Workshops, pp. 588-596.*
 Madhavan et al, "Semantic Mappings for Data Mediation", Feb. 2002, <pages.cs.wisc.edu/~anhai/talks/JayantMadhavan-Affiliates2002.ppt>, p. 1-17.
 Raghavan, S. et al. Crawling the Hidden Web. Computer Science Department, Stanford University, 2001, pp. 1-25, especially pp. 1-15.
 Don Clark, "AnswerFriend Seeks to Sell Question-Answering Software," Wall Street Journal, Aug. 24, 2000.
 Deniz Yuret, "Discovery of Linguistic Relations Using Lexical Attraction" PhD Thesis, MIT, May 15, 1998.
 Deniz Yurt "Lexical Attraction Models of Language." Submitted to The Sixteenth National Conference on Artificial Intelligence, 1999.
 Boris Katz, Deniz Yuret, et al. "Integrating Large Lexicons and Web Resources into a Natural Language Query System." To appear in Proceedings of the IEEE International Conference on Multimedia Computing and Systems, 1999.
 International Preliminary Report on Patentability and Written Opinion of the International Searching Authority for PCT/US2007/068597; International Searching Authority/US; Nov. 11, 2008.
 International Search Report for PCT/US2007/068597; International Searching Authority/US; May 22, 2008.
 International Preliminary Report on Patentability and Written Opinion of the International Searching Authority for PCT/US2007/075929; International Searching Authority/US; Feb. 17, 2009.
 International Search Report for PCT/US2007/075929; International Searching Authority/US; Jan. 21, 2008.

International Preliminary Report on Patentability and Written Opinion of the International Searching Authority for PCT/US2007/085646; International Searching Authority/US; Jun. 3, 2009.

International Search Report for PCT/US2007/085646; International Searching Authority/US; Jul. 23, 2008.

International Search Report for PCT/US2001/15711; International Searching Authority/US; Nov. 6, 2001.

International Preliminary Examination Report for PCT/US2001/15711; IPEA/US; Date of completion of report Feb. 16, 2005.

Stolowitz Ford Cowger LLP, Listing of Related Cases, Aug. 19, 2009. Internet Archive, Wayback Machine, website pages from LLBean, Feb. 29, 2000.

Chien, Steve, et al., "Semantic similarity between search engine queries using temporal correlation", In Proceedings of the 14th international conference on World Wide Web, Chiba, Japan; Session: Usage analysis, p. 2-11; published May 10-14, 2005; <http://www.ra.ethz.ch/CDstore/www2005/docs/p2.pdf>.

* cited by examiner

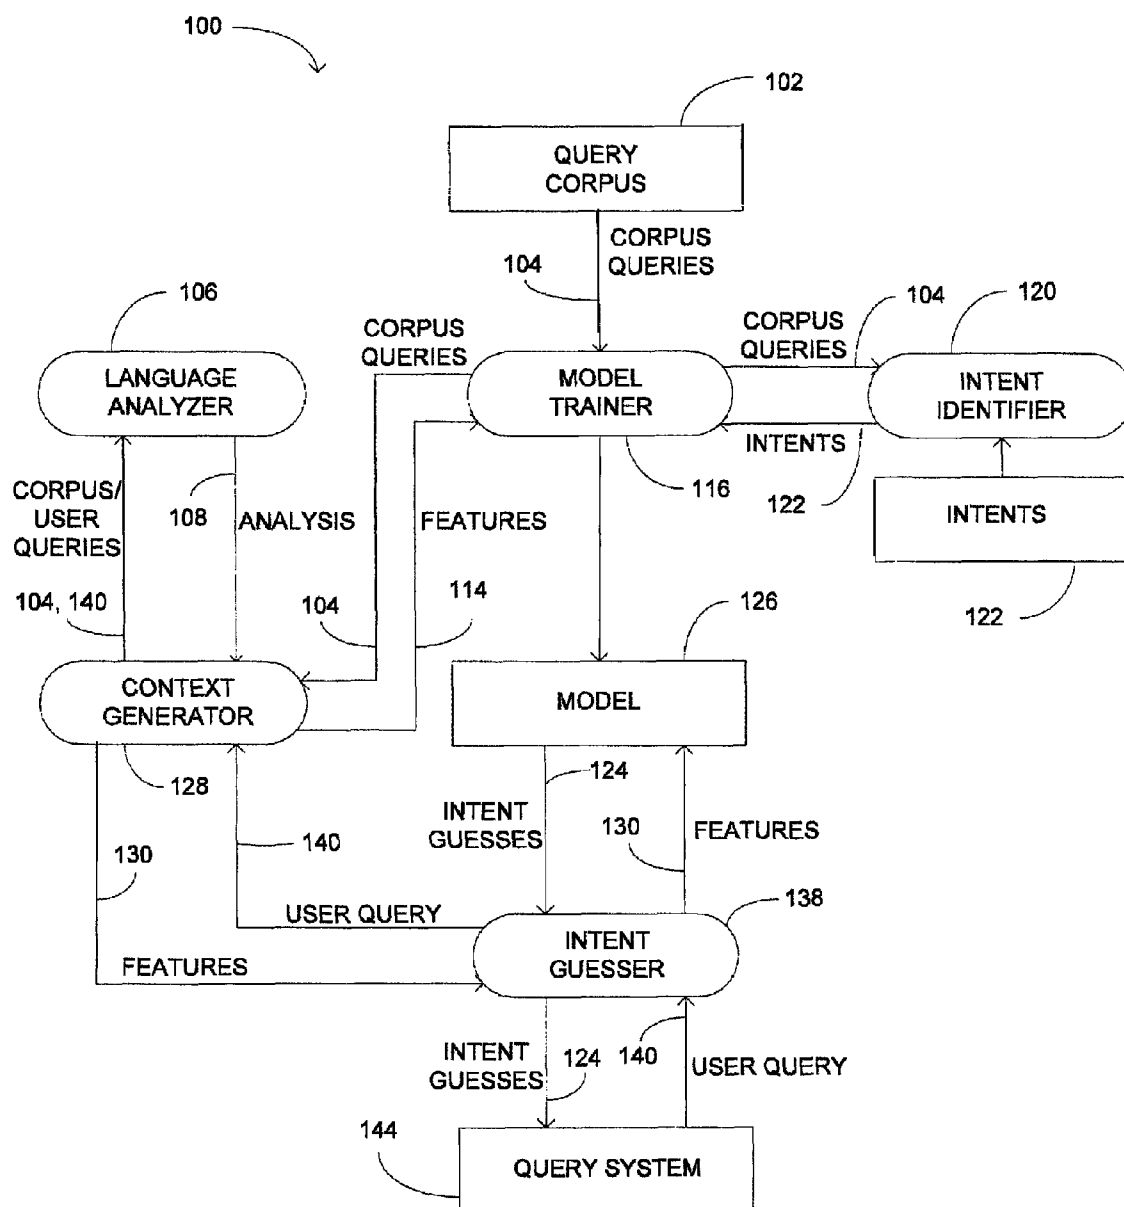


FIG. 1




FIG. 2

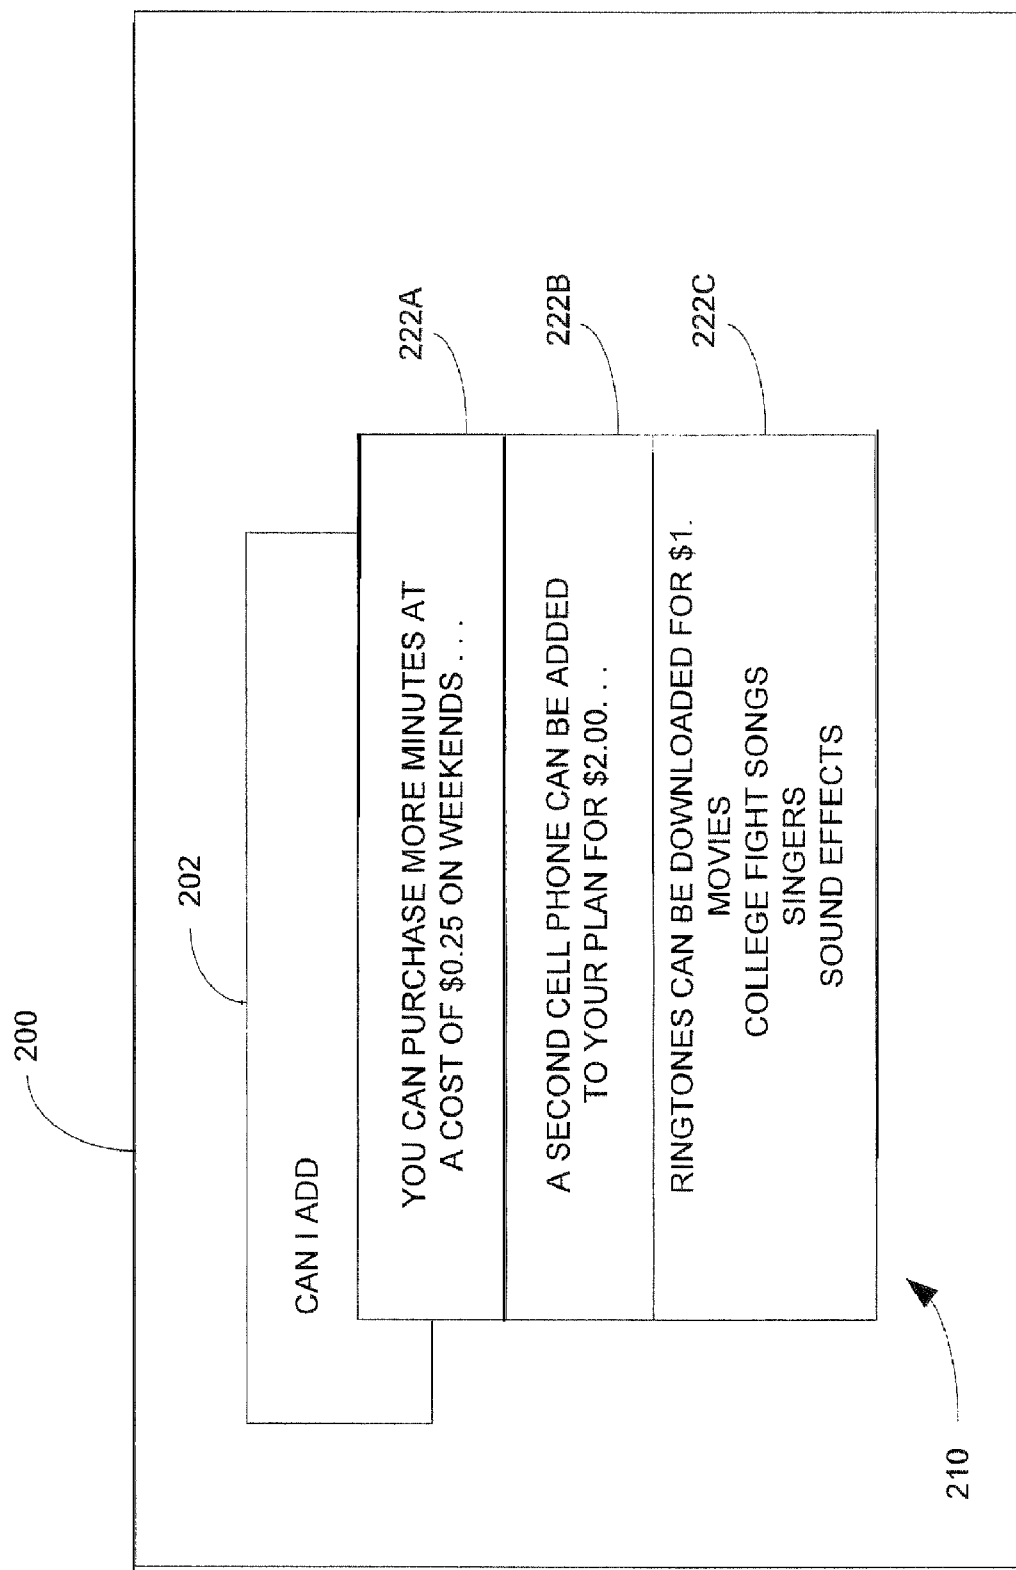


FIG. 3




FIG. 4




FIG. 5




FIG. 6

1

AUTOMATED SCHEME FOR IDENTIFYING USER INTENT IN REAL-TIME

The present application claims priority to provisional application Ser. No. 60/883,279, filed Jan. 3, 2007 which is incorporated by reference in its entirety.

The present application is also a continuation in part of U.S. patent application Ser. No. 10/820,341, filed Apr. 7, 2004, entitled: AN IMPROVED ONTOLOGY FOR USE WITH A SYSTEM, METHOD, AND COMPUTER READABLE MEDIUM FOR RETRIEVING INFORMATION AND RESPONSE TO A QUERY which is also incorporated by reference in its entirety and is a continuation in part of co-pending U.S. patent application Ser. No. 11/464,443, filed Aug. 14, 2006, entitled: METHOD AND APPARATUS FOR IDENTIFYING AND CLASSIFYING QUERY INTENT which is also herein incorporated by reference in its entirety.

BACKGROUND

Automatic information retrieval, search, and customer self service systems must, in some manner, understand an end user's query to a sufficient degree to be able to retrieve or construct an appropriate response. For keyword based systems this might simply involve the ability to extract keywords (important terms) from the query as well as interpret some simple operators such as NOT, OR, and PHRASE. For example, the following query: SHOW ME ALL DOCUMENTS CONTAINING THE WORDS "PYTHON" OR "COBRA" BUT NOT THE PHRASE "MONTY PYTHON" might be understood to mean: (PYTHON OR COBRA) NOT "MONTY PYTHON".

A more sophisticated system might understand the same meaning from the natural language query "SHOW ME INFORMATION ABOUT PYTHONS AND COBRAS" by understanding from the question's concepts that the desired answers are about snakes and not about the comedy troupe Monty Python. An even more sophisticated system might understand large classes of questions that effectively mean the same thing. For example, "HOW DO YOU CHECK YOUR BALANCE?", "HOW MUCH MONEY DO I HAVE?", "I'D LIKE TO VIEW MY CURRENT STATEMENT.", etc. may all be interpreted to mean the same thing. These question classes are called intents.

Some query systems attempt to understand a query while a user is completing the question. For example, a search engine may try to automatically provide suggestions for a search field of a web page while the user is still typing in the entry. This potentially reduces the time and effort required by the user to enter their query and possibly helps the user be more precise and avoid mistakes. This technique is primarily embodied as variants of what is often referred to as field "auto-completion". The system analyzes the query as the user is typing it in and proposes possible completions for the query from which the user can choose. For example, if the user types "golden" then the system might respond with "golden mean," "golden retriever," "Golden Gate," etc. These suggestions may be produced in any number of ways such as by rule based systems or statistical methods. However, all the suggestions begin with (or in some cases contain) the same text already input by the user.

BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 is a flow chart illustrating a system for guessing query intents.

2

FIG. 2 is a diagram showing how suggested intents and their responses might be presented.

FIG. 3 is a diagram for an alternative way to show suggested intents and their responses.

FIG. 4 is a diagram showing how parameterized intents might be handled.

FIG. 5 is a diagram showing how weighted intents might be handled.

FIG. 6 is a diagram describing an environment for the intent guessing system in FIG. 1.

DETAILED DESCRIPTION

Intent Guesser

FIG. 1 illustrates an intent guessing system **100** that goes beyond simple auto-completion to guessing user intents as the user is typing a query. This frees the system **100** from providing only guesses that match the same text as the user's current input. Intent guessing uses a statistical classification system to produce a model **126** which can then be used by an intent guesser **138** to provide possible intent guesses. The system **100** in one example uses a maximum entropy classification system, however other embodiments may use other classification systems that may produce comparable results. A classification system associates inputs with an outcome via features created from the input. This process results in a model which can then produce possible outcomes (with probabilities indicating their likely correctness) given a new input. In the case of intent guesser **138**, the input is a user query **140**, the outcome is an intent guess **124**, and the features **130** used to produce the outcome **124** based on a linguistic analysis of the user query **140**.

To explain in more detail, a model trainer **116** is software that creates a model **126** from a query corpus **102**. The query corpus **102** is a list of questions similar to the questions that might be entered into the query system **144**. For example, an enterprise server may continuously log queries that are entered into an enterprise web-site. The queries are stored in memory and all or some subset of the past queries may be used in query corpus **102**.

This query corpus **102** may be updated from time to time to reflect the most recent and/or most common questions that are being asked by people accessing the enterprise website. For example, an updated query corpus **102** may contain common questions that enterprise customers ask about new enterprise products. The query corpus **102** need not be an exhaustive list of all possible questions (as this is not possible), but the larger the corpus, the better the results are likely to be.

A model trainer **116** forwards the corpus queries **104** to a context generator **128** and receives back features **114**. The context generator **128** creates the features **114** from a linguistic analysis of the corpus queries **104** using a language analyzer **106** which is described in co-pending U.S. patent application Ser. No. 10/820,341, filed Apr. 7, 2004, entitled: AN IMPROVED ONTOLOGY FOR USE WITH A SYSTEM, METHOD, AND COMPUTER READABLE MEDIUM FOR RETRIEVING INFORMATION AND RESPONSE TO A QUERY which is incorporated by reference in its entirety.

The context generator **128** feeds the corpus queries **104** to the language analyzer **106** and receives back analysis data **108** that identifies the different language characteristics. For example, the context generator **128** in combination with the language analyzer identifies different words and different concepts in the corpus queries **104**. The context generator **128** sends these different features **114** back to the model trainer **116**. Examples of features **114** include, but are not limited to: query tokens (e.g., words, numbers, punctuation marks); their

stems (the words stripped of morphology—"dog" instead of "dogs"); and concepts (e.g. <canine> instead of "dog" or "mutt").

The model trainer 116 sends the same corpus queries 104 to an intent identifier 120. The intent identifier 120 identifies the intents 124 that match the corpus queries 104 and sends the matching intents 122 back to the model trainer 116. Generating intents 122 and matching intents with queries is described in co-pending U.S. patent application Ser. No. 11/464,443, filed Aug. 14, 2006, entitled: METHOD AND APPARATUS FOR IDENTIFYING AND CLASSIFYING QUERY INTENT which is herein incorporated by reference in its entirety.

The intents 122 may be created by an enterprise specialist based on the information provided on the enterprise website, the subject matter associated with the enterprise and industry, and the questions previously submitted to the website by different people.

The model trainer 116 receives the features 114 back from the context generator 128 and receives the intents 122 from intent identifier 120 that are all associated with those same corpus queries 104. The model trainer 116 creates a model 126 that affiliates the corpus query features 114 with different intents 122. The exact nature of the model 126 will vary depending on the technology used, but in one embodiment comprises a data structure with statistical associations between features and outcomes. There are many implementation possibilities for of the model 126 that are known to those knowledgeable in the field.

Query Intent Guessing

The query system 114 uses the intent guesser 138 to retrieve a list of intent guesses 124 ordered by their probable correctness. Like the model trainer 116, the intent guesser 138 submits user queries 140 received from query system 144 to the context generator 128. Similar to the corpus queries 104, the context generator 128 uses the language analyzer 106 to decompose the user queries 140 into different features 130 that may include, but are not limited to, tokens (e.g., words, numbers, punctuation marks); their stems (the words stripped of morphology—"dog" instead of "dogs"); and concepts (e.g., <canine> instead of "dog" or "mutt").

The intent guesser 138 feeds the features 130 received from the context generator 128 into model 126 and receives back a list of one or more intent guesses 124. The intent guesses 124 are provided to the query system and then displayed to the user that submitted the original user query 140.

It is worth noting that intents 124 are pieces of information that attempt to surmise the thought or "intent" associated with the query 140 while the user is entering the query and possibly before the user has completed the entire query 140. Accordingly, the intent guesses 124 may not necessarily include any of the words used in the corpus queries 104 or any of the actual words used in the user query 140. This makes the intent guessing system 100 extremely powerful, since the actual words entered by the user in query 140 do not necessarily restrict what responses can be provided by the intent guessing system 100.

This is particularly pertinent for query guesses 124 that are provided while the query is still being entered. Initial words used in a query may have little or no relevance for conventional auto-fill-in software. However, in an intent guessing system, the first few words could provide substantial information about what the user is actually thinking. Thus, the intent guessing system 100 would be able to identify and complete the intended relevant query for the user more quickly.

The intent guesses 124 can also provide more relevant query responses. For example, a user may not know the precise name of a product. Common rule based query systems may never be able to accurately complete or respond to the user query 140, since the user never provides the appropriate key words. However, the intent guessing system 100 has the capacity to surmise what the user is thinking based on the meanings of the words and concepts in the query 140 and the pre-created intents 124 associate with those words and concepts. The system 100 can derive this user intent even when the user does not use the correct key words in the query 140.

As an example with respect to cell phones, if the user were to type "CAN I ADD" as user query 140, the intent guesser system 100 might suggest intent guesses 124 of: "ADDITIONAL LINE OR PHONE", "PURCHASING MINUTES", etc. As another example, typing the user query 140 "WHAT IS THE RATE" might produce intent guesses 124 of: "SERVICE PLAN MINUTES RESEARCH", "INTERNATIONAL CALLING RESEARCH", AND "SERVICE PLAN RESEARCH". However, if the user completes the query 140 to "WHAT IS THE RATE FOR CALLING AUSTRALIA?", the system 100 might produce a more precise intent guess 140: "INTERNATIONAL CALLING RESEARCH." It should be noted from these examples that this is quite distinct from auto-completion where such a leap from initial portions of a thought to a final thought intent are not possible.

Presentation

FIG. 2 shows a computer 190 for accessing and displaying a web page 200 that serves as an interface to the query system 100 shown in FIG. 1. FIG. 2 shows one way in which intent guesses 204A-204C may be displayed to an end user, as a list of choices in a drop down menu 210 below a query box 202. The query system 144 accesses intent guessing system 100 (FIG. 1) in response to a user entering the partial user query 140 "CAN I ADD". The intent guessing system 100 then returns the intent guesses 204A-204C shown in drop down menu 210 as described above in FIG. 1.

In addition, as described in the co-pending U.S. patent application Ser. No. 11/464,443, filed Aug. 14, 2006, entitled: METHOD AND APPARATUS FOR IDENTIFYING AND CLASSIFYING QUERY INTENT, any combination of intents 204A-204C may be associated with intent responses, such as intent response 206, which is meant to be presented to the end user when an associated intent guess matches the query 140.

The responses 206 may be promotions, links, or any other helpful piece of information relevant to the intent 204. These responses 206 may be displayed within the framework of the intent guesser system 100, for example, by showing the intent response 206 when the end user hovers a cursor over the associated intent guess 204A-204C. This has the benefit of allowing the user to see more information about the intent guesses 204A-204C before selecting one, and also allows the user to take a desired action without ever issuing the query to a search system.

Example

As an example, referring to FIGS. 1 and 2, query corpus 102 may contain many earlier logged queries including the following possibilities:

1. Adding minutes to my plan
2. Can I include more minutes?
3. Is it possible to add an additional phone line?
4. How do I put ringtones on my phone?

The model trainer 116 is responsible for generating a model 126 from these queries. The model trainer 116 uses the context generator 128 and language analyzer 106 to identify relevant features 114 from the corpus queries 104. For

5

example, the features **114** for query (1) might include the tokens “adding”, “minutes”, “to”, “my”, and “plan”; the stems “add”, “minute”, “to”, “my”, and “plan”; and the concepts <add>, <minute>, and <plan>. The features **114** for query (2) might include the tokens “can”, “i”, “include”, “more”, and “minutes”; the stems “can”, “i”, “include”, “more”, and “minute”; and the concepts <add> and <minute>. Notice that although the queries are quite different, they share many of the same features **114**, for example the token “minutes” and the concept <add>.

As described in the aforementioned co-pending patent application, these queries are associated with intents: queries (1) and (2) might be associated with the intent **204A** “PURCHASING MINUTES”; query (3) with intent **204B** “NEW LINE OR PHONE”; and query (4) with the intent **204C** “DOWNLOADING RINGTONES”.

The model trainer **116** generates a model **126** using the features **114** and the intents **204** mentioned above. Again as mentioned above, the exact nature of the model **126** will vary depending on the technology used, but may be some data structure with statistical associations between features and outcomes.

At this point, the query system **144** might receive a user query **140** “CAN I ADD”, as shown in FIG. 2. The intent guesser **138** uses the context generator **128** and language analyzer **106**, just like the model trainer **116**, to determine the features **130** of the partial user query “CAN I ADD”. The features **130** may include the tokens “can”, “i”, and “add”; the stems “can”, “i”, and “add”; and the concept <add>. These features **130** are then sent to the model **126** which outputs the intent guesses **124** most likely to be associated with features **130**.

In the example shown in FIG. 2, the model **126** determines that the intent guesses **204A-204C**: “PURCHASING MINUTES”, “NEW LINE OR PHONE”, and “DOWNLOADING RINGTONES”, respectively, are all possibilities. The intent **204A** “PURCHASING MINUTES” is listed first as the most likely intent because the model **126** determined that intent **204A** is most closely associated with the features in query **140** than the other intents **204B** and **204C**. For example, intent **204A** has more instances of the concept <add> plus the tokens “can” and “i”. The resulting intents **204** are presented to the user as shown in FIG. 2.

FIG. 3 shows an alternative embodiment for displaying the intent responses **222A-222C**. Instead of displaying intent **204A** in FIG. 2, the query system **144** may display a response **222A** associated with the intent **204A**. In this case, the response **222A** is a link to a web page that allows a user to purchase additional cell phone minutes. Similarly, instead of displaying intent **204B** in FIG. 2, the query system **144** may display an intent response **222B** that when selected links the user to a web page for purchasing a new phone line. Any combination of intents **204** in FIG. 2 and intent responses **222** in FIG. 3 can be displayed on the web page **200** as a drop down menu **210** or as separate icons **206**.

Parameterized Intents

FIG. 4 shows an example of a parameterized intent **124A** that is sent back by the model **126**. The parameterized intent **124A** may include a generalized concept, such as the concept **125** <PRODUCT>. The model **126** may send back a parameterized intent **124A** when the identified intent may require a specific type of information from the query **140**. For example, the model **126** may determine that the query **140** is asking for help about a product. However, there may not be a particular intent associated with the product referred to in the user query **140**. Accordingly, the model sends back a generalized and

6

parameterized intent **124A** that identifies the concept <PRODUCT> and the word HELP.

The concept **125** in intent **124A** causes the intent guesser **138** to search through the associated user query **140** for one or more words that can replace the <PRODUCT> concept **125**. In this example, the intent guesser **138** identifies the words CELL PHONE in query **140** that correspond to the concept <PRODUCT> in intent **124A**. The intent guesser **138** replaces the concept <PRODUCT> with the words CELL PHONE and sends the revised intent **124B** to the query system **144**.

The query system **144** then displays the revised intent **124B** CELL PHONE HELP to the user. If the intent CELL PHONE HELP is selected by the user, then the query system **144** sends the intent to a search engine **252** to retrieve the related information.

Intent Probability

FIG. 5 shows a set of intents **150** that are output by the model **126** in response to the features **130** for a particular user query. Each intent **150A-150E** may have an associated probability value **154** that indicates the probability of the intent correctly matching the intent of the user query. The probabilities **154** can be assigned by the model **126** according to the number of features **130** in the user query **140** that match features associated with the intents **150**. Of course other parameters can also contribute to the probability values **154**, such as matches for particular types of word or concepts.

In order to provide only the most relevant intents to the user, the intent guesser **138** uses one or more thresholds to determine which intents are sent back to the query system **144**. It may happen that many intents fall within different ranges where a first high probability range contains most of the intents likely to be responsive to the user query **140**. Other intents may be banded within lower ranges that are much less likely to provide helpful responses.

Accordingly, the intent guesser **138** only forwards the intents **150** in the upper probability range to the query system **144**. For example, in FIG. 5, the two intents **150A** and **150B** are both in an upper probably range and all other intents **150C-150E** are in substantially lower ranges. Accordingly, only intents **150A** and **150B** are sent back to query system **144** in response to user query **140**.

Environment

The operations described in FIGS. 1-4 are performed by a computer, processor, and/or computer server. Furthermore, the system described above can use dedicated processor systems, micro controllers, programmable logic devices, or microprocessors that perform some or all of the operations. Some of the operations described above may be implemented in software and other operations may be implemented in hardware.

For example, FIG. 6 shows a computer network system **240** that includes an enterprise network **248** that has one or more enterprise servers **250** and one or more enterprise databases **254** that contain content associated with enterprise network **248**. For example, the enterprise network **248** may be operated by a cell phone service provider. In this example, the enterprise database **16** may contain information related to cell phone purchasing and service. In another example, the enterprise network **248** may be associated with a car manufacturer or car dealership and the enterprise database **254** includes vehicle information. These are, of course, just two examples and any other type of business or entity can also be associated with enterprise network **248**.

Other web servers **270** may operate outside of the enterprise network **248** and may include associated web files or other web content **280**. Examples of content stored in enter-

prise database **254** and in web file server **270** may include HTML web pages, PDF files, Word® documents, structured database information or any other type of electronic content that can contain essentially any type of information.

Some of the information may be stored in a structured format referred to generally as structured content. Data may be stored in the enterprise database **254** in a preconfigured format specified for enterprise network **248**. For example, a cell phone service price list may be considered structured content. Alternatively, other information that is contained in enterprise database **254**, or contained on other web server **270**, may be considered non-structured content. This may include HTML web pages, text documents, or any other type of free flowing text or data that is not organized in a preconfigured data format known by the query system **144**.

A query may be initiated from the computer **190** through a User Interface (UI) **200** that in one example may be a web page displayed by a web browser. The computer **190** in one example may be a Personal Computer (PC), laptop computer, wireless Personal Digital Assistant (PDA), cellular telephone, or any other wired or wireless device that can access and display content over a packet switched network **260**. In this example, the query **140** is initiated from the UI **200** and transported over the Internet **260** to the enterprise server **250**.

The novel intent guessing system **100** provides one or more intent guesses **204** to the query system **144** as described above. The intent guesses **204** provided by the intent guessing system **144**, along with other responses, are then sent back to computer **190** for displaying on the UI **200**.

The user may select any of the identified intents **204**. The selected intent is sent back to the enterprise server **250**. The query system **144** then sends the selected intent **204** to search engine **252** to search for associated structured and/or non-structured content either in the enterprise database **254** or web server **270**. Any results from the search are then sent back to the terminal **200** via the query system **144**. In another embodiment, some or all of the query system **144**, intent guessing system **100** and search engine **252** may be located and operated on terminal **190**.

The enterprise server **250** may include one or more processors that are configured to operate the query system **144**, intent guessing system **100**, and search engine **252**. The operations performed by the server **250** could be provided by software computer instructions that are stored in a computer readable medium, such as memory on server **250**. The instructions are then executed by a processor in server **250**. It should be understood that the examples presented below are used for illustrative purposes only and the scope of the invention is not limited to any of the specific examples described below.

For the sake of convenience, the operations are described as various interconnected functional blocks or distinct software modules. This is not necessary, however, and there may be cases where these functional blocks or modules are equivalently aggregated into a single logic device, program or operation with unclear boundaries. In any event, the functional blocks and software modules or features of the flexible interface can be implemented by themselves, or in combination with other operations in either hardware or software.

Having described and illustrated the principles of the invention in a preferred embodiment thereof, it should be apparent that the invention may be modified in arrangement and detail without departing from such principles. Claim is made to all modifications and variation coming within the spirit and scope of the following claims.

The invention claimed is:

1. A method of processing a current user query, comprising:
 - training a model, said training comprising:
 - continuously logging previous user queries;
 - extracting a question from the continuously logged previous user queries;
 - analyzing the extracted question to obtain at least one feature;
 - assigning an intent to the extracted question; and
 - adding an entry to the model and mapping the obtained feature for the extracted question to the assigned intent for the extracted question in the added entry of the model;
 - receiving an initial portion of the current user query, wherein the initial portion includes a particular string of text;
 - guessing an intent associated with the initial portion of the current user query while a remainder of the current user query is being received, wherein said guessing using only the initial portion further comprises:
 - analyzing the initial portion to obtain at least one feature;
 - comparing, using a processing device, the obtained at least one feature of the initial portion to entries in the model;
 - selecting, using the processing device, at least one entry from the model based on the comparison; and
 - determining the intent for the initial portion of the current user query according to an intent included in the selected entry;
 - wherein the intent guess obtained for the initial portion according to said analyzing, said comparing, said selecting, and said determining, does not include any portion of the particular text string; and
 - responding to the initial portion of the current user query with the intent guess.
2. The method according to claim 1, wherein the at least one feature includes at least one selected from the group comprising query tokens, query token stems, and concepts.
3. The method according to claim 1, further comprising:
 - identifying a parameterized intent guess that includes at least one concept;
 - identifying words in the current user query that are associated with the concept;
 - modifying the parameterized intent guess by replacing the concept with the identified words from the current user query; and
 - responding to the current user query with the modified parameterized intent guess.
4. The method according to claim 1, further comprising:
 - identifying a response associated with the intent guess; and
 - either displaying the response along with the associated intent guess or displaying the response instead of the associated intent guess.
5. The method according to claim 4, wherein the response includes links to web pages or additional information associated with the intent guess.
6. The method according to claim 1, further comprising:
 - identifying a plurality of intent guesses associated with the current user query;
 - identifying an order for the intent guesses of the plurality according to how well the intent guesses match the current user query; and
 - displaying the intent guesses in the identified order.
7. An article of manufacture including a computer-readable medium having instructions stored thereon that, in response

9

to execution by a computing device, cause the computing device to perform operations comprising:

receiving an initial portion of a current user query, wherein the initial portion includes a particular string of text;
 training a model, said training comprising:
 continuously logging previous user queries;
 extracting a question from the continuously logged previous user queries;
 analyzing the extracted question to obtain at least one feature;
 assigning an intent to the extracted question; and
 adding an entry in the model and mapping the obtained feature for the extracted question to the assigned intent for the extracted question in the added entry of the model;
 guessing an intent associated with the initial portion of the current user query while a remainder of the current user query is being received, wherein said guessing using only the initial portion further comprises:
 inputting the initial portion into a context generator to obtain at least one feature corresponding to the initial portion;
 analyzing the obtained feature of the initial portion according to the model, wherein the analysis includes comparing the obtained at least one feature of the initial portion to entries in the model;
 selecting, using the processing device, at least one entry from the model based on the comparison; and
 determining the intent for the initial portion of the current user query according to an intent included in the selected entry;
 wherein the intent guess obtained for the initial portion according to said analysis, said selecting, and said determining does not include any portion of the particular text string; and
 responding to the initial portion of the current user query with the intent guess.

8. The article of manufacture according to claim 7, wherein the operations further comprise:

identifying a first subset of words from the initial portion while the current user query is being entered;
 making an intent guess based on the first subset of words; and
 replying to the first subset of words with the intent guess while the current user query is still being entered.

9. The article of manufacture according to claim 7, wherein the operations further comprise:

identifying at least one question class that includes at least one concept;
 identifying words in the current user query that are associated with the concept;
 modifying the identified question class by replacing the concept with the identified words in the current user query; and
 responding to the current user query with the modified question class.

10. The article of manufacture according to claim 9, wherein the operations further comprise:

identifying confidence factors for the question classes; and
 displaying the question classes according to the identified confidence factors.

10

11. The article of manufacture according to claim 9, wherein the question classes include associated probability values that indicate a probability that the question classes correctly respond to the current user query and the question classes are sent or not sent in response to the current user query according to the associated probability values.

12. The article of manufacture according to claim 7, wherein the operations further comprise:

identifying additional information associated with the intent guess; and
 either displaying the identified additional information with the intent guess or displaying the additional information instead of the intent guess.

13. An apparatus, comprising:

a processor; and
 a memory to store instructions that, when executed by the processor, are operable to:

receive an initial portion of a current user query, wherein the initial portion includes a particular string of text;
 train a model, said training comprising:

continuously logging previous user queries;
 extracting a question from the continuously logged previous user queries;
 analyzing the extracted question to obtain at least one feature;

assigning an intent to the extracted question; and
 adding an entry to the model and mapping the obtained feature for the extracted question to the assigned intent for the extracted question in the added entry of the model;

guess an intent associated with the initial portion of the current user query while a remainder of the current user query is being received, wherein the guess using only the initial portion includes the following operations:

extracting at least one feature from the initial portion;
 comparing, using the processing device, the extracted at least one feature to entries from the model;
 selecting, using the processing device, at least one entry from the model based on the comparison; and
 determining the intent for the initial portion of the current user query according to an intent included in the selected entry;

wherein the intent guess obtained for the initial portion according to said extracting, said comparing, said selecting, and said determining does not include any portion of the particular text string; and
 respond to the initial portion of the current user query with the intent guess.

14. The apparatus of claim 13, wherein the processing device is configured to:

identify a parameterized intent guess that includes at least one concept;
 identify words in the current user query that are associated with the concept;
 modify the parameterized intent guess by replacing the concept with the identified words from the current user query; and
 respond to the current user query with the modified parameterized intent guess.

15. The apparatus of claim 13, wherein the at least one feature includes at least one selected from the group comprising query tokens, query token stems, and concepts.

* * * * *